

LLMs for AI Planning.

Harsha Kokel

<https://harshakokel.com>





James Oswald
RPI



Daniel Cao
Cornell University



Kavitha Srinivas
IBM Research



Michael Katz
IBM Research



Shirin Sohrabi
IBM Research

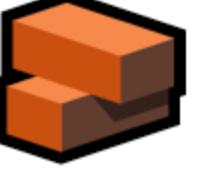


Outline

- LLMs as ~~3D Printers~~ Mold Makers
- LLMs as MoldMakers for AI Planning
 - LLM for search code generation
 - LLM for NL2PDDL
- Benchmarking and Evaluations

Mold Makers vs 3D Printers

Use cases?

-  Make a mold once and replicate fast, is consistent, and scalable
-  Costly and slow, used for bespoke, unique objects

LLMs as 3D Printers

Generate each piece from scratch

- Engineer Every Response
 -  Time and token costs
 -  Inconsistent format, tone, and logic
 -  Redundant efforts

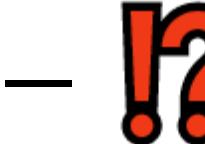
LLMs as Mold Makers

Build Once, Use Many Times

Use the LLM to:

- Create reusable templates, frameworks, or data structures
- Generate code, logic flows, or prompts as "molds"

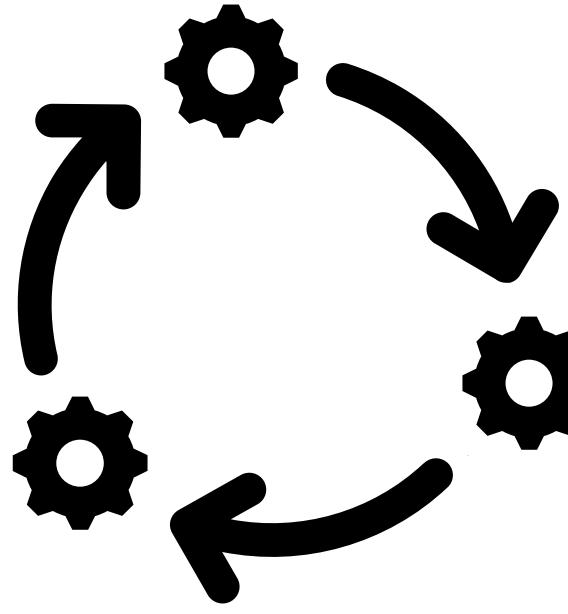
Benefits:

-  Faster inference
-  Lower costs
-  Consistency and reliability  

Outline

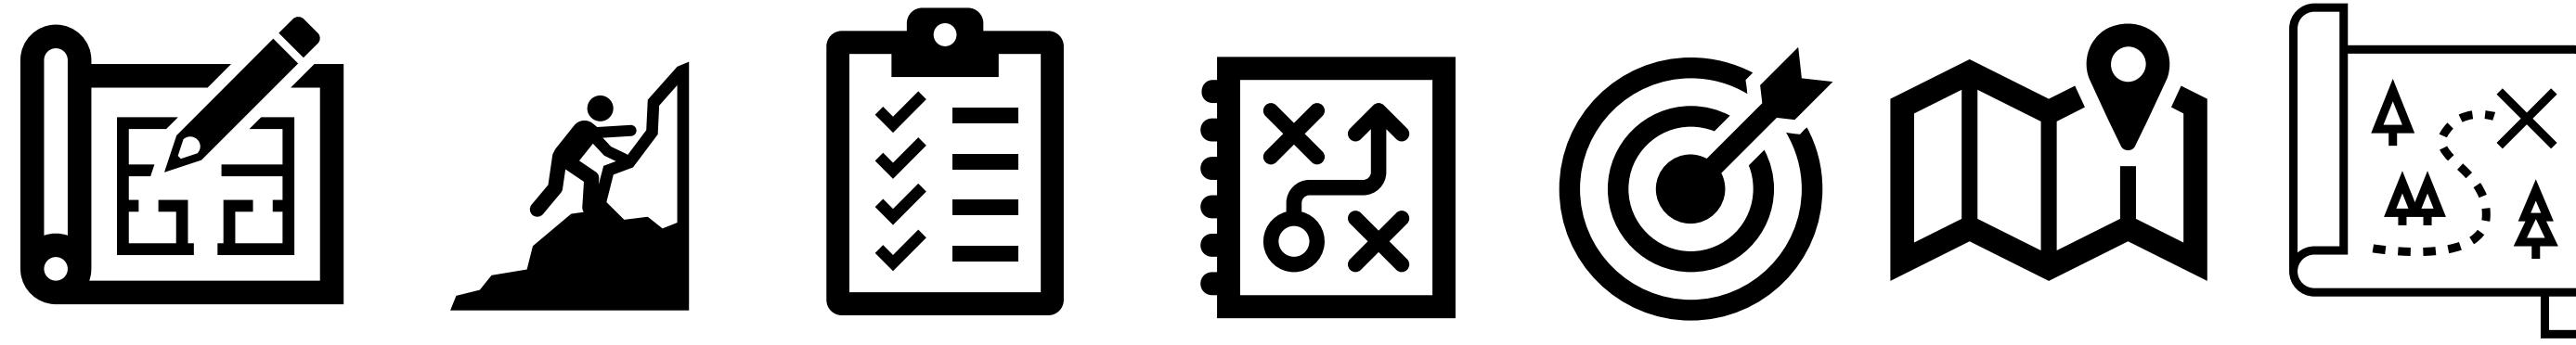
- LLMs as ~~3D Printers~~ Mold Makers
- LLMs as MoldMakers for AI Planning
 - LLM for search code generation
 - LLM for NL2PDDL
- Benchmarking and Evaluations

What is AI Planning?



AI Planning is a sub-field of AI
that explores autonomous techniques to solve
planning problems.

What is a Planning Problem?



Given the following description

- an initial state
- a goal (or objective)
- a set of actions that transform the state

devise a sequence of action that achieves the goal

LLMs as Planners



LLMs as 3D Printers

COT: Reasoning Abilities

Standard Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The answer is 27. 

Chain-of-Thought Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

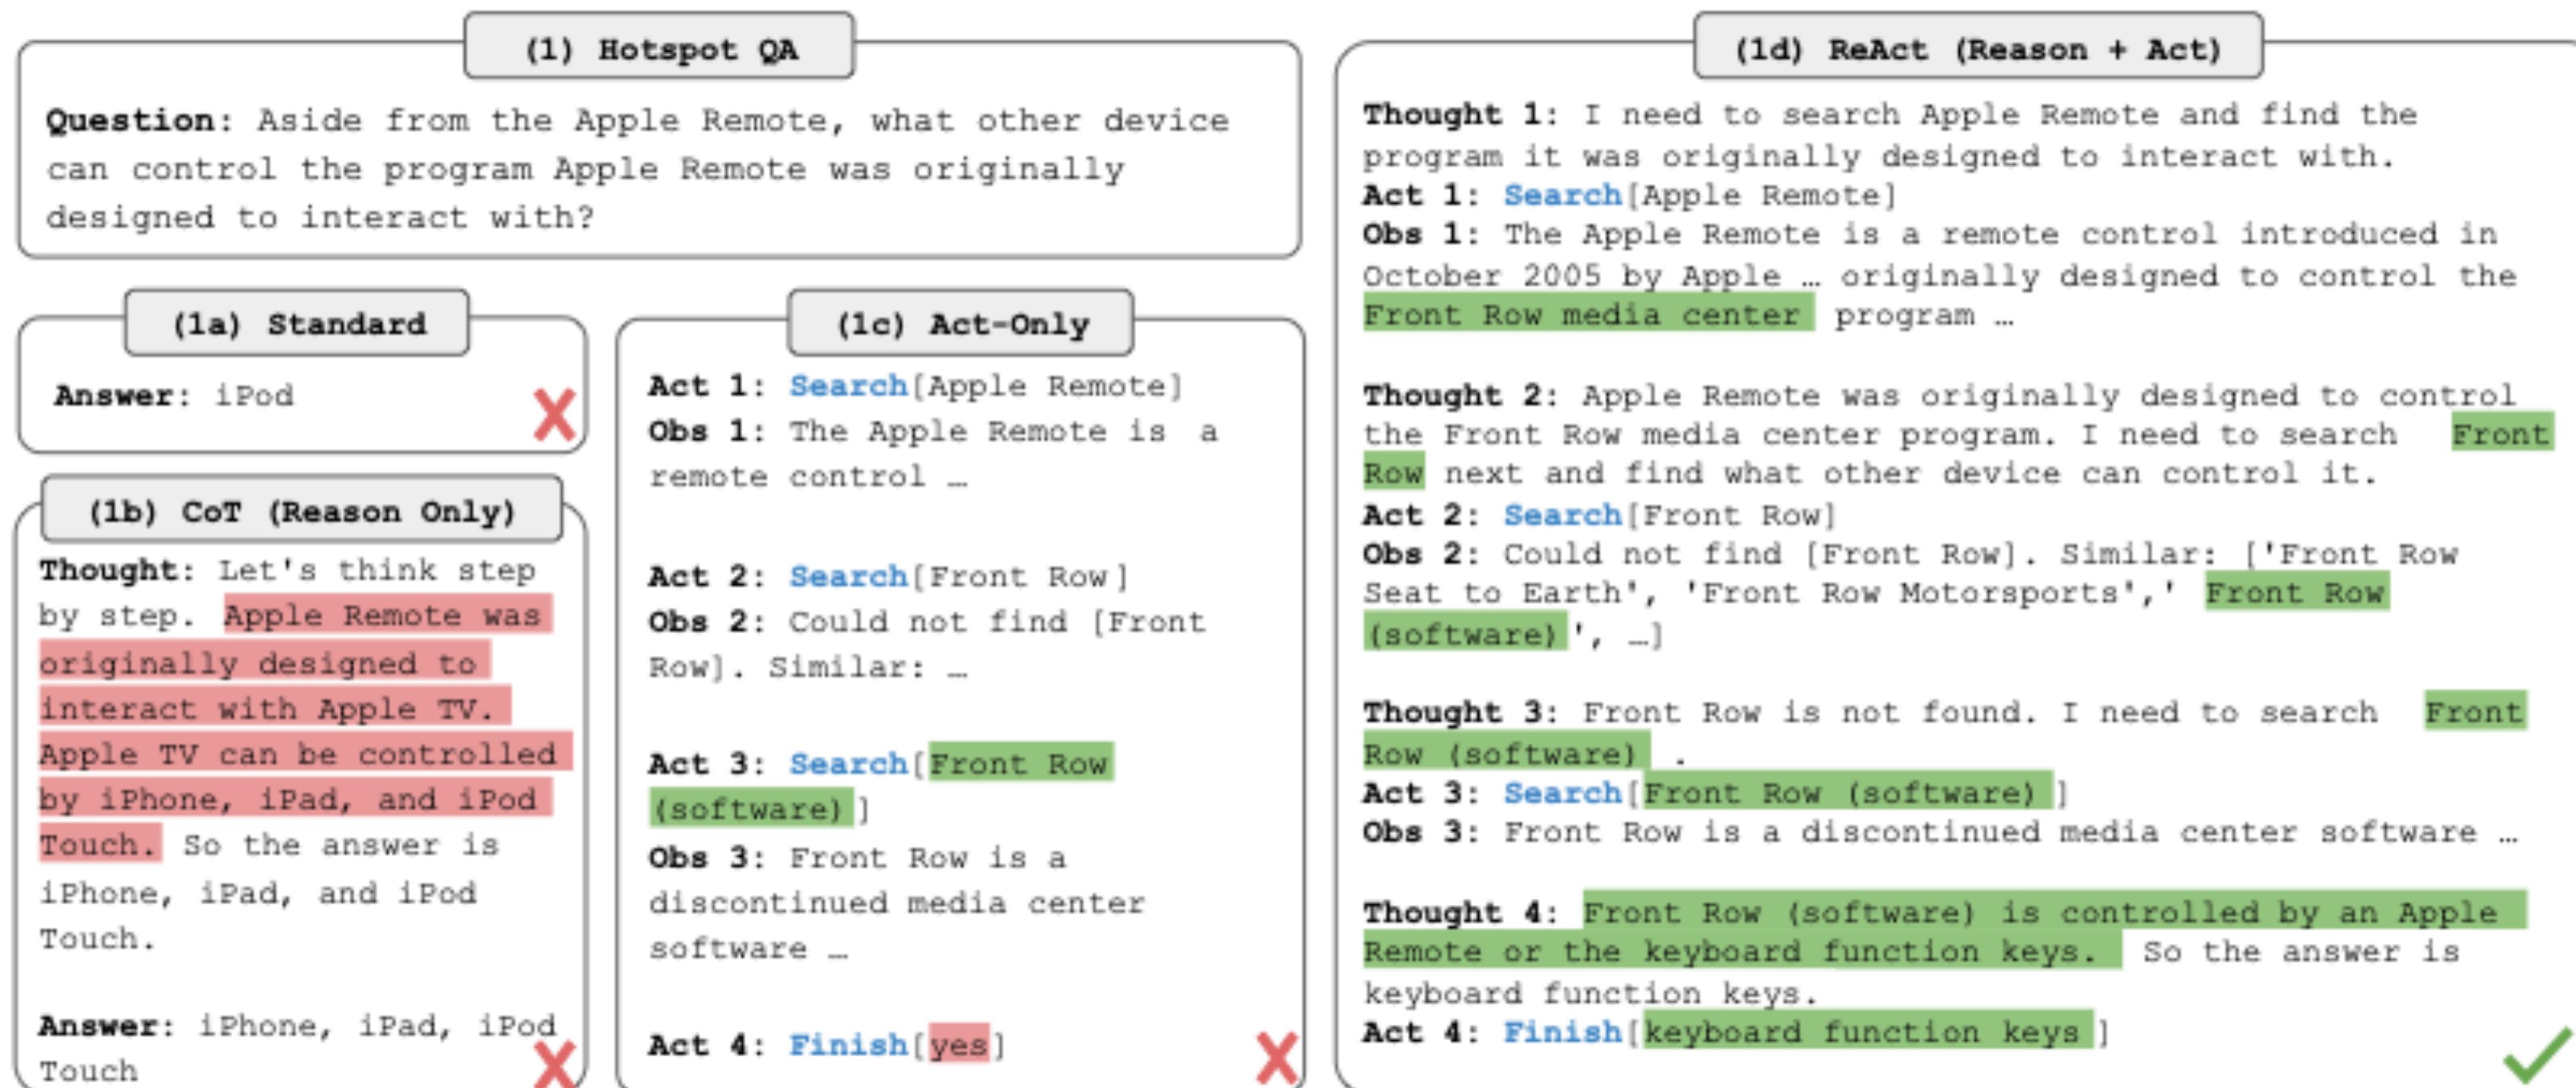
Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

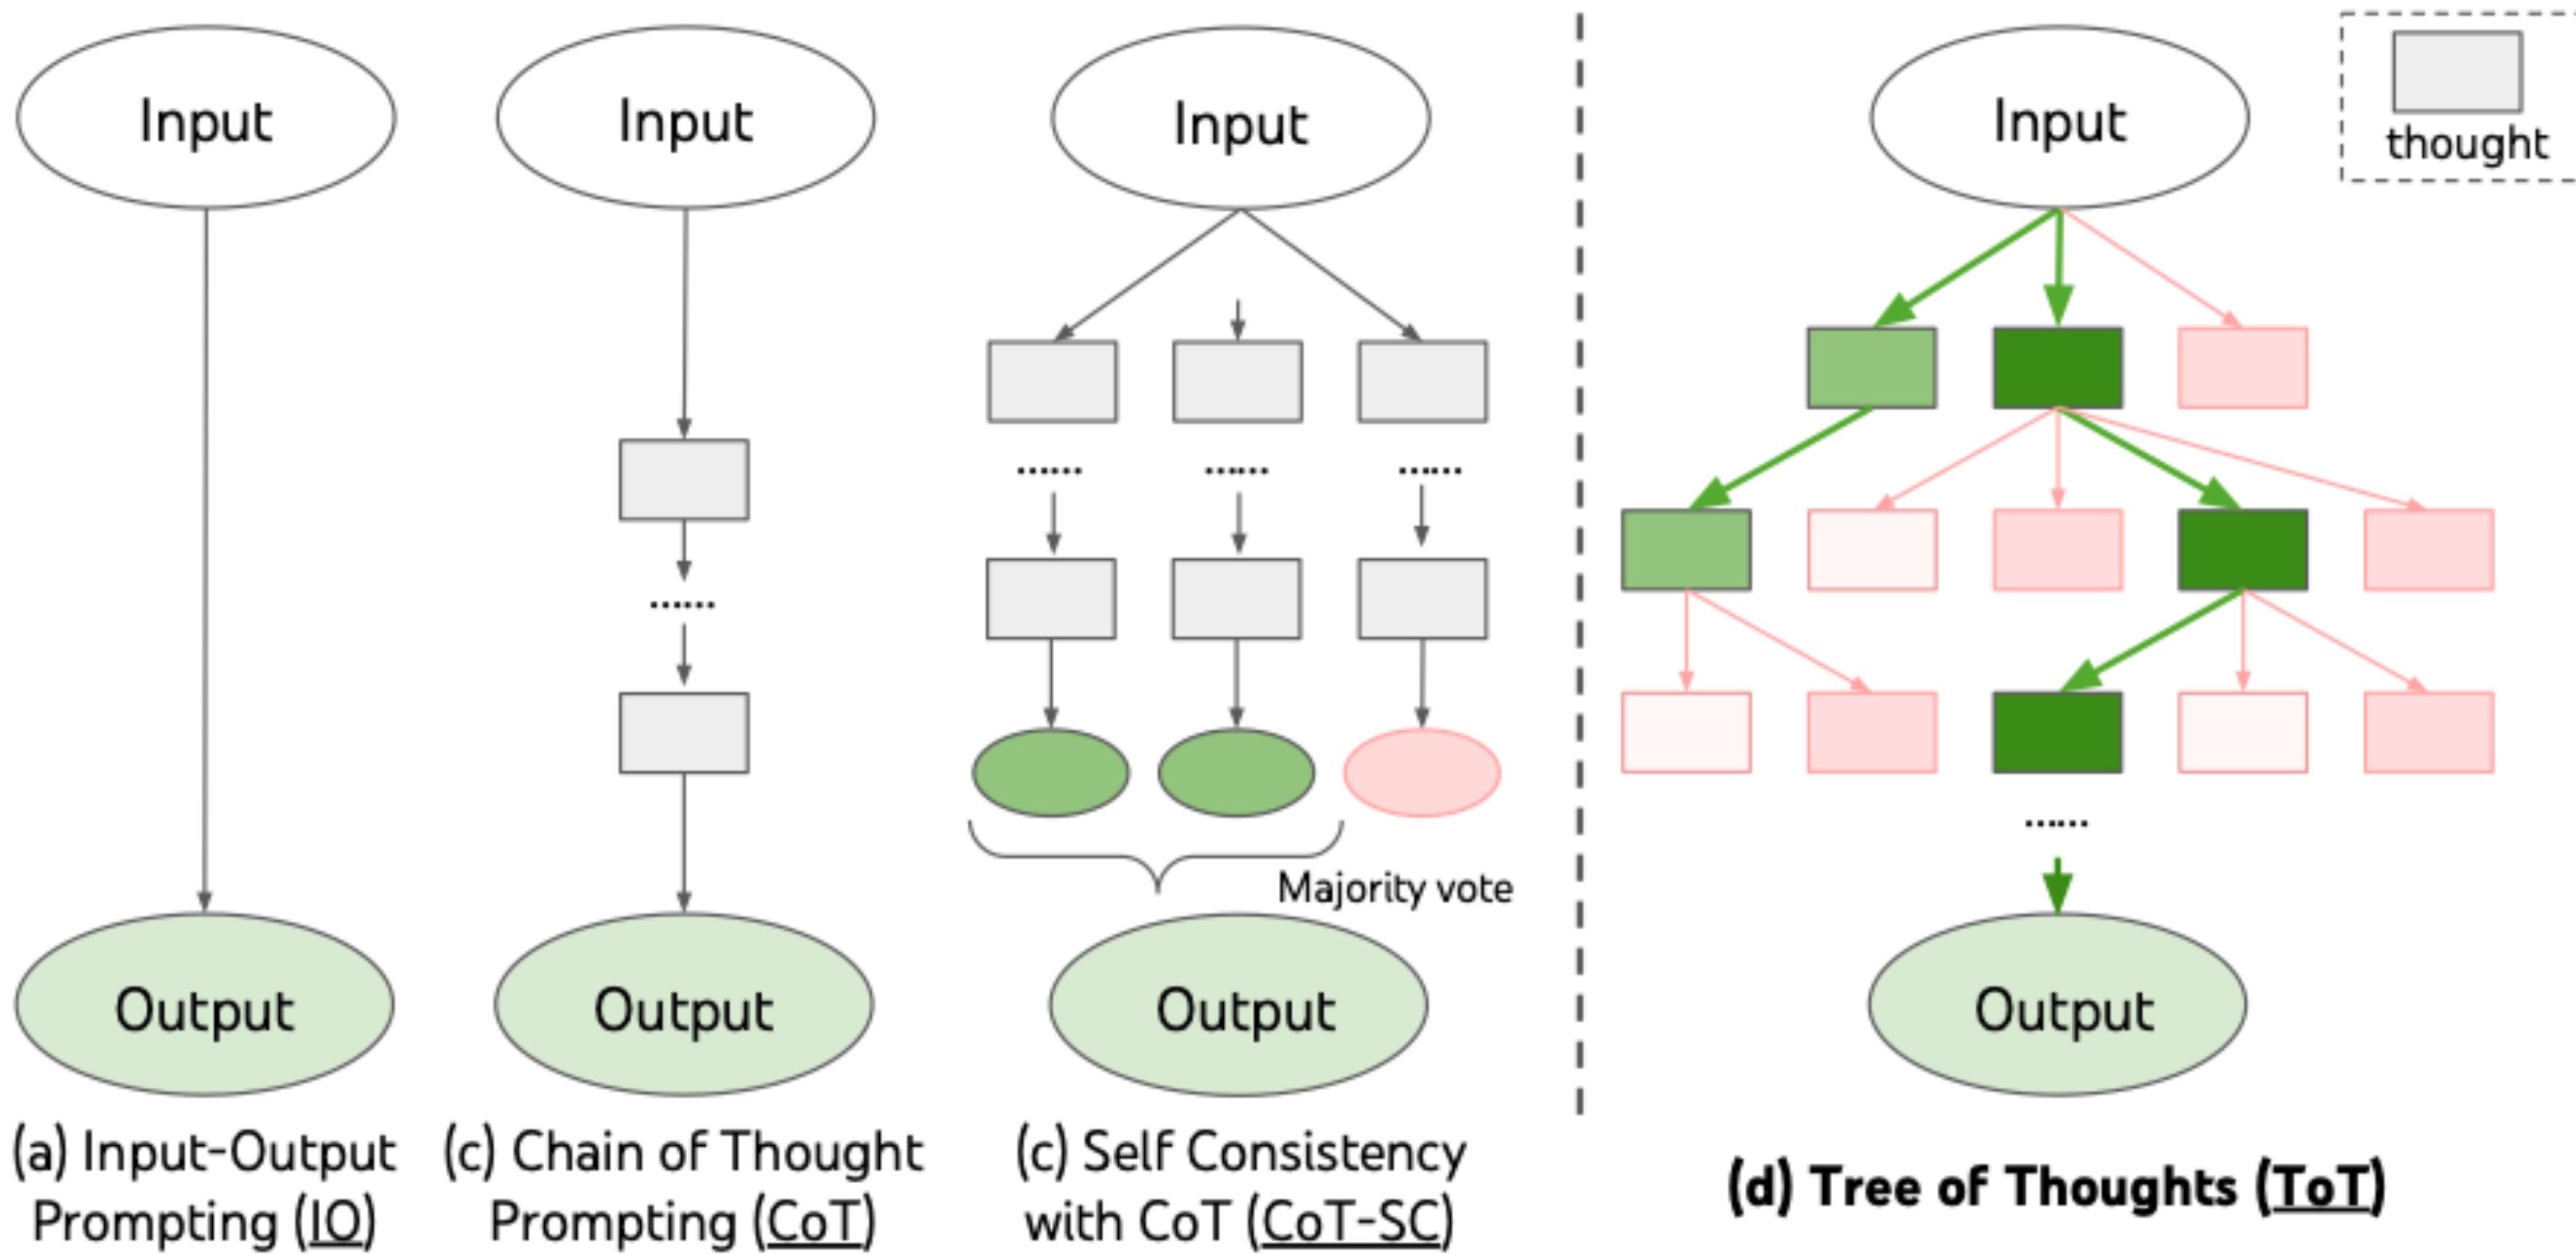
A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9. 

ReAct:

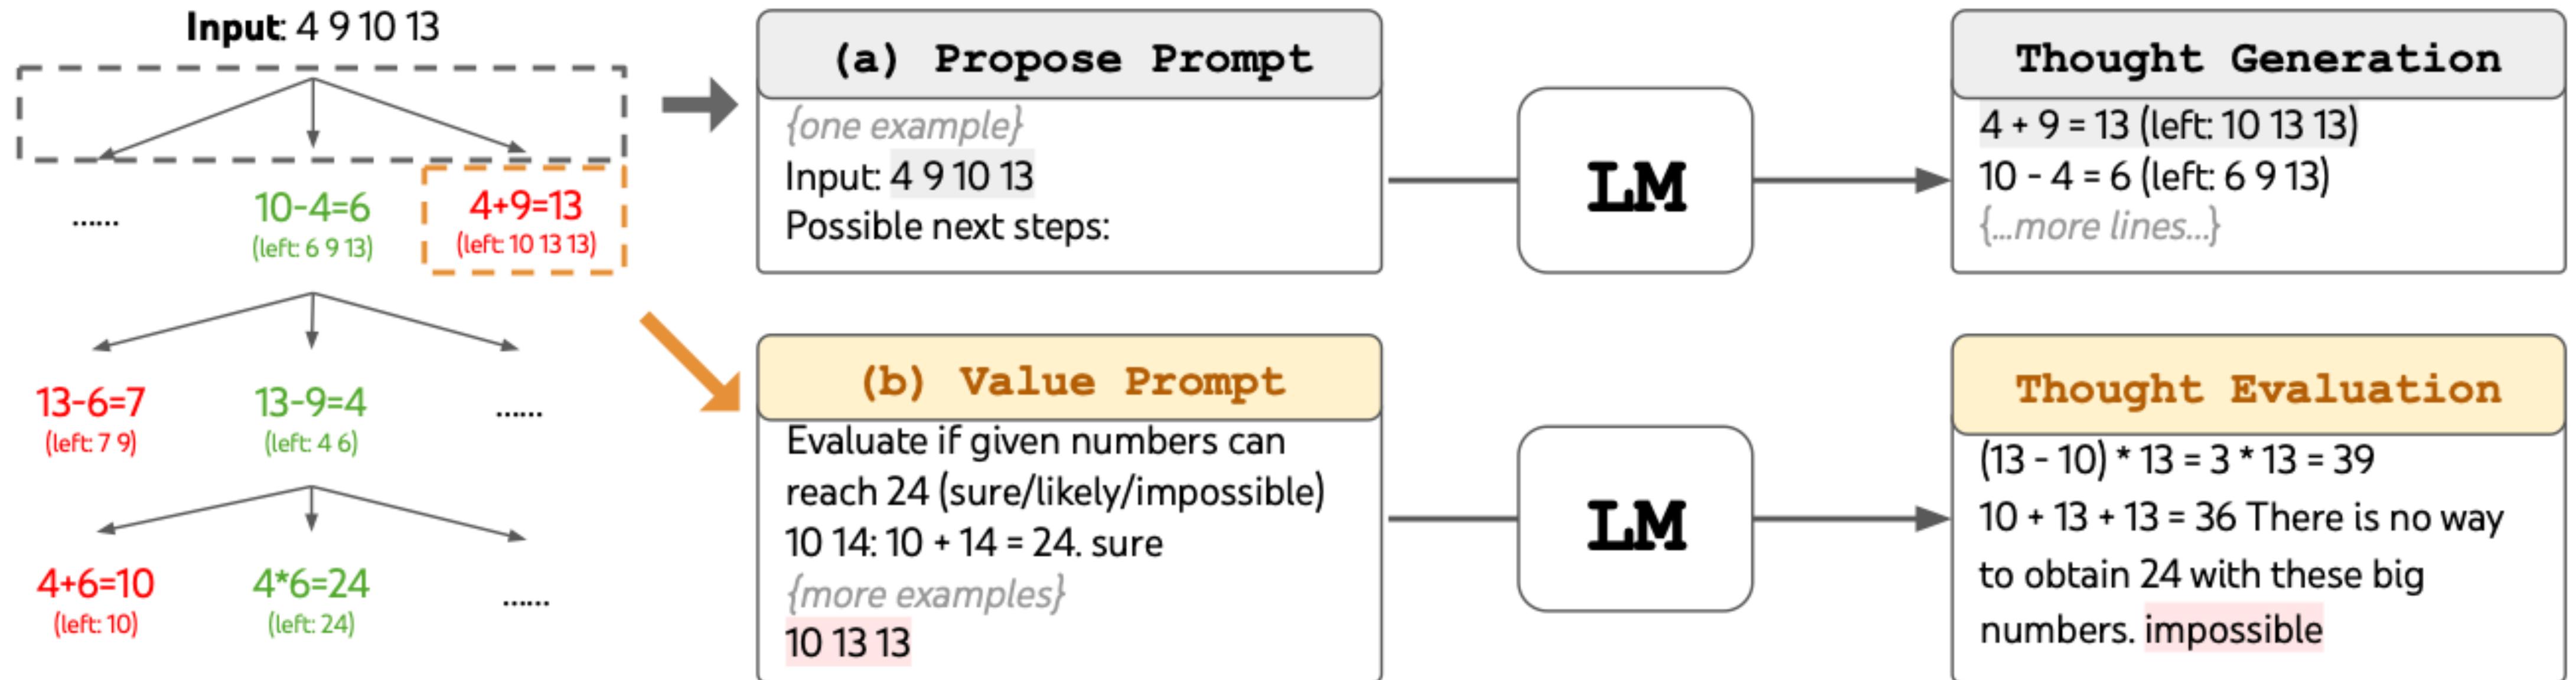
Reasoning and Tool Use/Acting



Tree of Thoughts: Search



Game of 24



Calls: 1

Chain-Of-Thoughts **Reflexion**
IO ReWOOCT
ReAct Tree-of-thoughts
ToT RAP Graph-of-thoughts
Reasoning-via-Planning **LATS**

and many more...

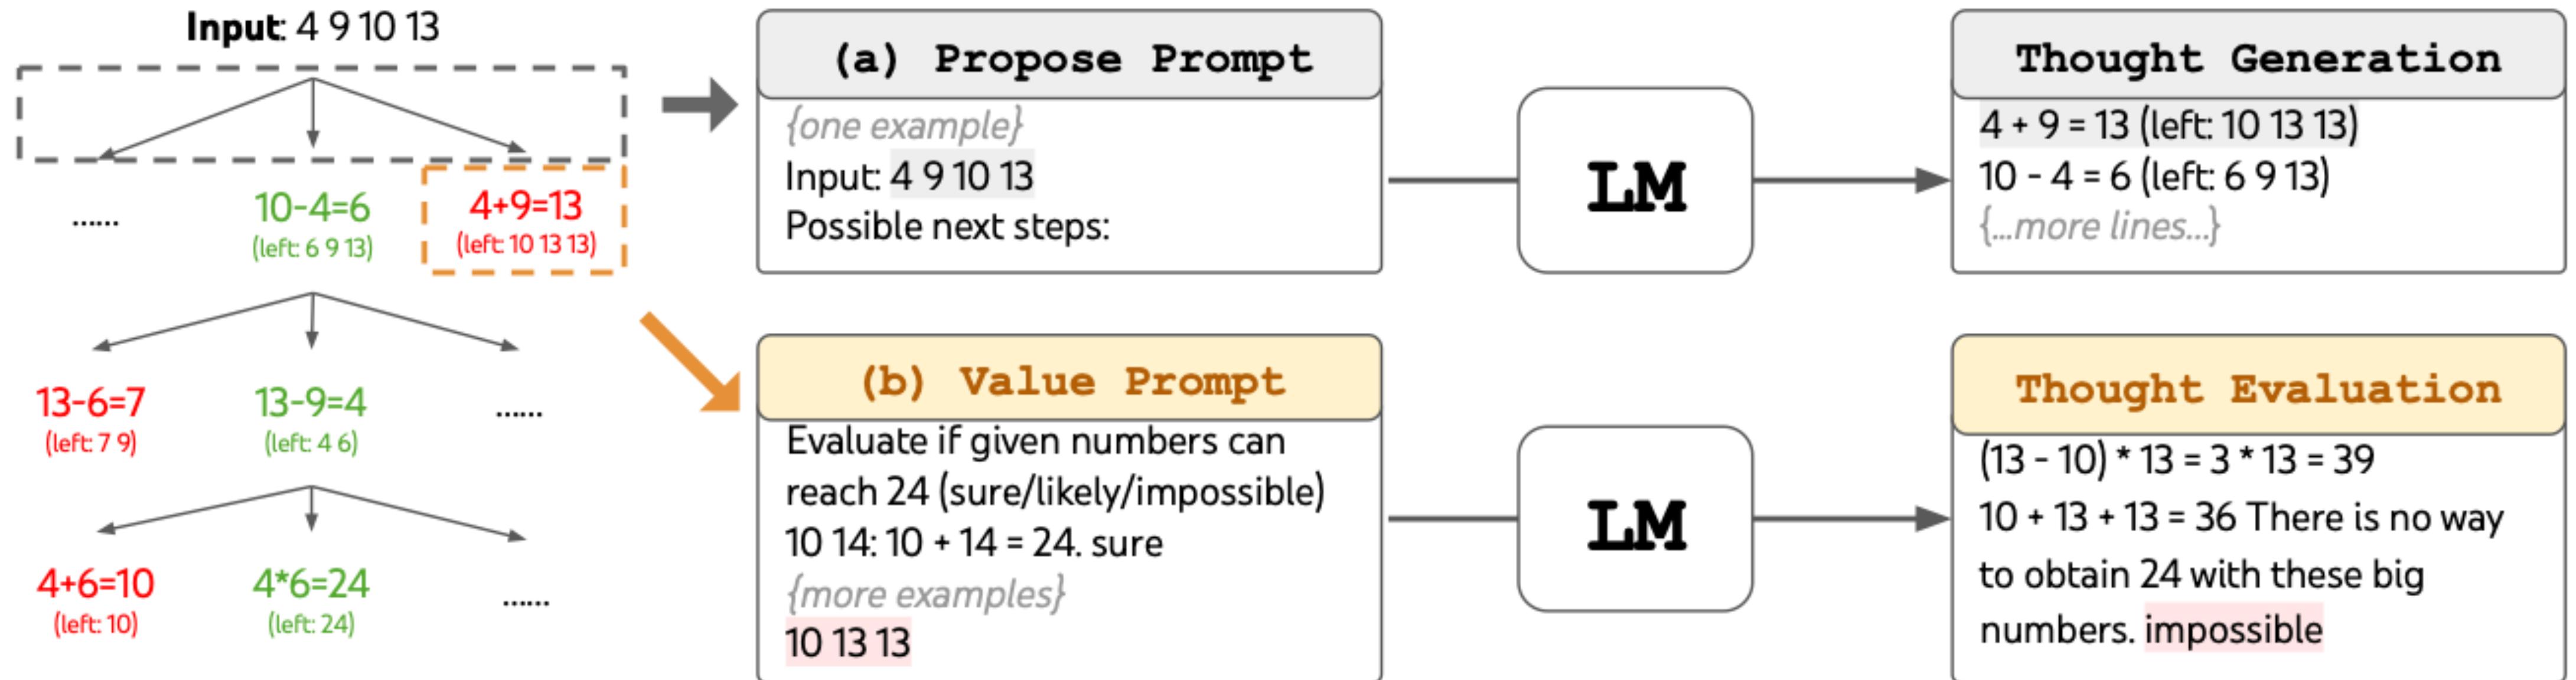
- No guarantees of soundness
 - o solution generated may or may not work
- No guarantees of completeness
 - o can miss correct solutions
- No efficiency considerations
 - o extremely high # LM Evaluation
 - o tackle one problem at a time
 - o worse !! Generate one action at a time

They abandon soundness and completeness for the sake of inefficiency !

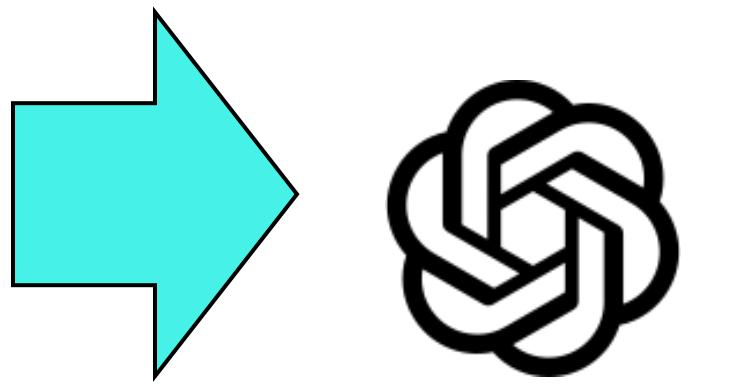
Outline

- LLMs as ~~3D Printers~~ Mold Makers
- LLMs as MoldMakers for AI Planning
 - **LLM for search code generation**
 - LLM for NL2PDDL
- Benchmarking and Evaluations

Game of 24



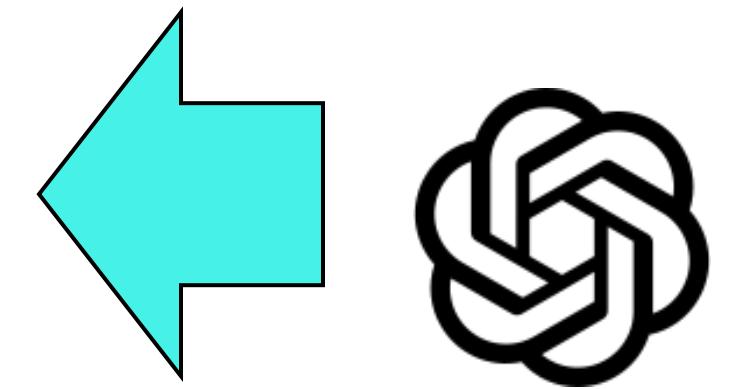
User: The 24 Game is a mathematical card game in which the objective is to find a way to manipulate four integers so that the end result is 24. The game is played with a list of four numbers, and the player must use all four numbers exactly once, using any combination of addition, subtraction, multiplication, or division, to arrive at the number 24. If a state is defined by a list of numbers (4 or less), how are successor states defined? Provide only the python code that returns a list of successor states for a state.



ChatGPT

Below is the Python code that defines how to generate the successor states for a given state in the 24 Game. A state in the 24 Game is a list of numbers. The successor states are generated by taking two numbers from the state, applying an operation (addition, subtraction, multiplication, or division if the divisor isn't zero), and then forming a new list with the result and the remaining unused numbers.

```
import itertools
def generate_successor_states(state):
    successors = []
    n = len(state)
    for (i, j) in itertools.combinations(range(n), 2):
        a, b = state[i], state[j]
        remaining = [state[k] for k in range(n) if k != i and k != j]
        # Generate all possible new states
        successors.append(remaining + [a + b])
        successors.append(remaining + [a * b])
        if a > b:
            successors.append(remaining + [a - b])
        else:
            successors.append(remaining + [b - a])
        if b != 0:
            successors.append(remaining + [a / b])
        if a != 0:
            successors.append(remaining + [b / a])
    return successors
```



ChatGPT

User: The 24 Game is a mathematical card game in which the objective is to find a way to manipulate four integers so that the end result is 24. The game is played with a list of four numbers, and the player must use all four numbers exactly once, using any combination of addition, subtraction, multiplication, or division, to arrive at the number 24. If a state is defined by a list of numbers (4 or less), how are successor states defined? Provide only the python code that returns a list of successor states for a state.

ChatGPT: Below is the Python code that defines how to generate the successor states for a given state in the 24 Game. A state in the 24 Game is a list of numbers. The successor states are generated by taking two numbers from the state, applying an operation (addition, subtraction, multiplication, or division if the divisor isn't zero), and then forming a new list with the result and the remaining unused numbers.

```
import itertools

def generate_successor_states(state):
    successors = []
    n = len(state)
    for (i, j) in itertools.combinations(range(n), 2):
        a, b = state[i], state[j]
        remaining = [state[k] for k in range(n) if k != i and k != j]

        # Generate all possible new states
        successors.append(remaining + [a + b])
        successors.append(remaining + [a * b])
        if a > b:
            successors.append(remaining + [a - b])
        else:
            successors.append(remaining + [b - a])
        if b != 0:
            successors.append(remaining + [a / b])
        if a != 0:
            successors.append(remaining + [b / a])

    return successors
```

This function takes a state (a list of numbers), iterates over all pairs of numbers, applies the operations to these pairs, and then forms new lists with the result of the operation and the other numbers that were not used in the operation. This ensures that each successor state is a valid step towards finding a solution to reach 24, using all four numbers exactly once along with any of the allowed operations.

User: Provide only the python code that tests whether a state is a goal state.

ChatGPT: Here's the Python code that tests whether a given state is a goal state in the 24 Game. A goal state is achieved if the list contains exactly one number, and that number is 24.

```
def is_goal_state(state):
    return len(state) == 1 and state[0] == 24
```

This function checks if the length of the state list is 1 and whether the sole element in the list is the number 24.

Thought of Search

- Human can verify the code
- Plug it in to any search algorithm
 - o guaranteed correct solution
- Efficient
 - o For game of 24, all 1362 games in < 2 sec
 - o 1.2 queries to LLM on average for sound successor function

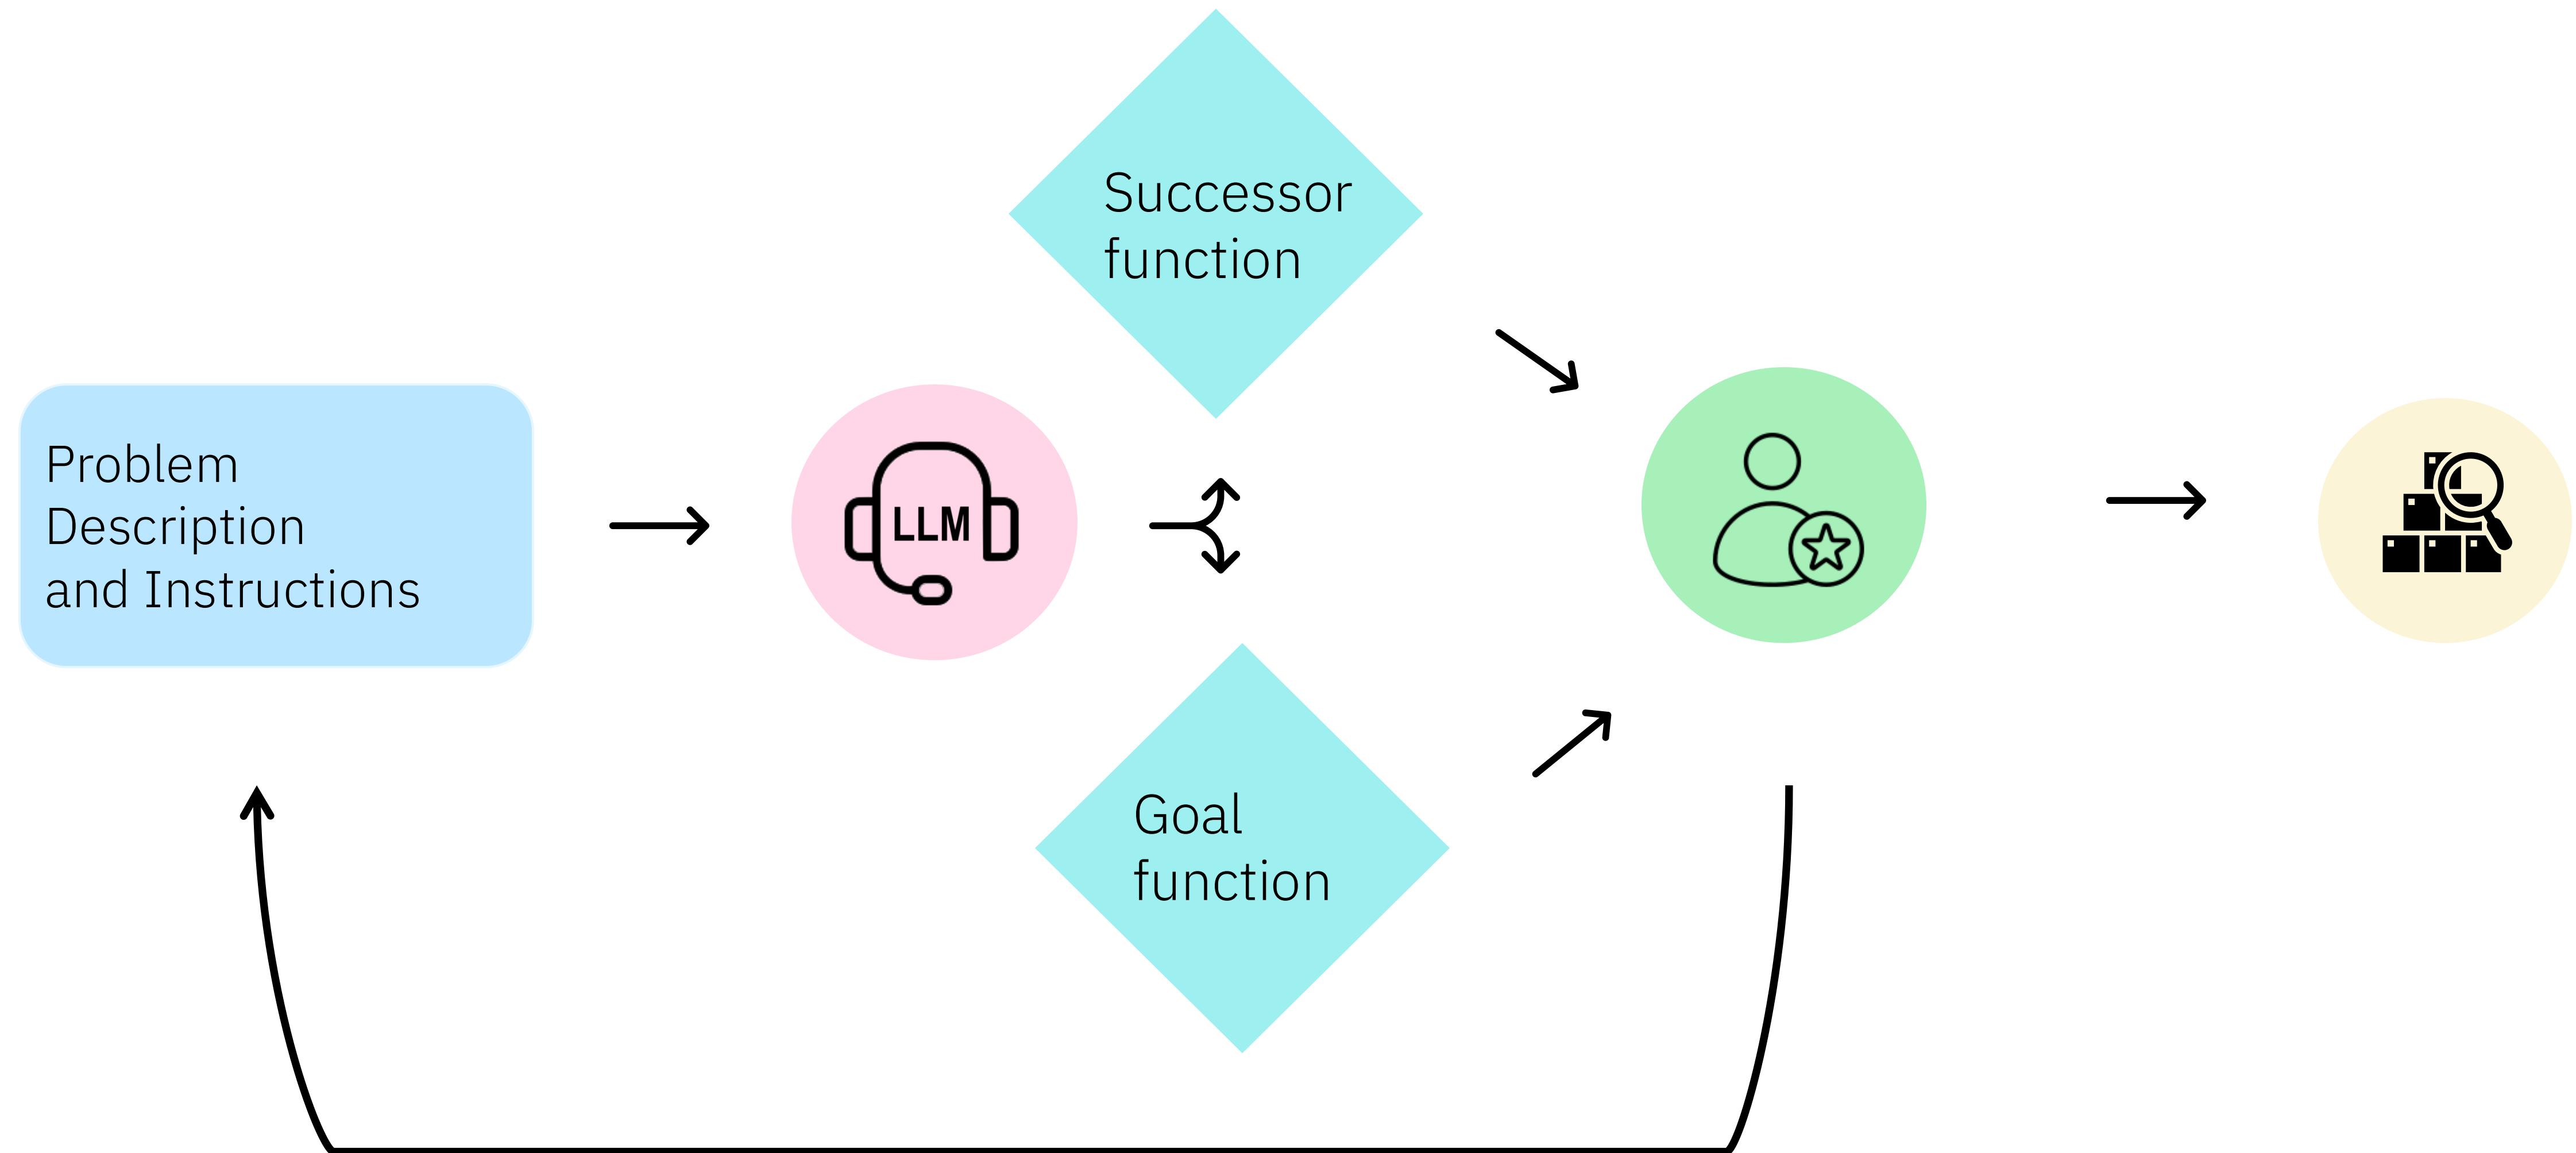
Experiments

Approach	Complexity	24Game		Crossword		BlocksWorld		PrOntoQA	
		States	Calls	States	Calls	States	Calls	States	Calls
IO	$O(D)$	0.02%	1362	4e-9%	20	0.5%	502	4%	4000
CoT	$O(D)$	0.02%	1362	4e-9%	20	0.5%	502	4%	4000
ReAct	$O(LD)$	0.07%	4086	4e-8%	200	7.8%	8032	24.6%	24K
ReWOO	$O(LD)$	0.07%	4086	4e-8%	200	7.8%	8032	24.6%	24K
RAP	$O(TbLD)$	3.3%	245K	2e-6%	12K	388%	482K	1229%	1.44M
ToT	$O(bmLD)$	1.6%	102K	1e-6%	5K	194%	201K	615%	600K
GoT	$O(bLD)$	0.3%	20K	2e-7%	1K	39%	40K	122%	120K
Reflection	$O(LTD)$	0.7%	68K	4e-7%	2.4K	77.6%	90K	245%	320K
LATS	$O(TbLD)$	3.3%	286K	2e-6%	14K	388%	562K	1229%	1.68M
ToS (ours)	$O(1)$	27.0%	2.2	3e-4%	3.8	125%	3.8	175%	2.6

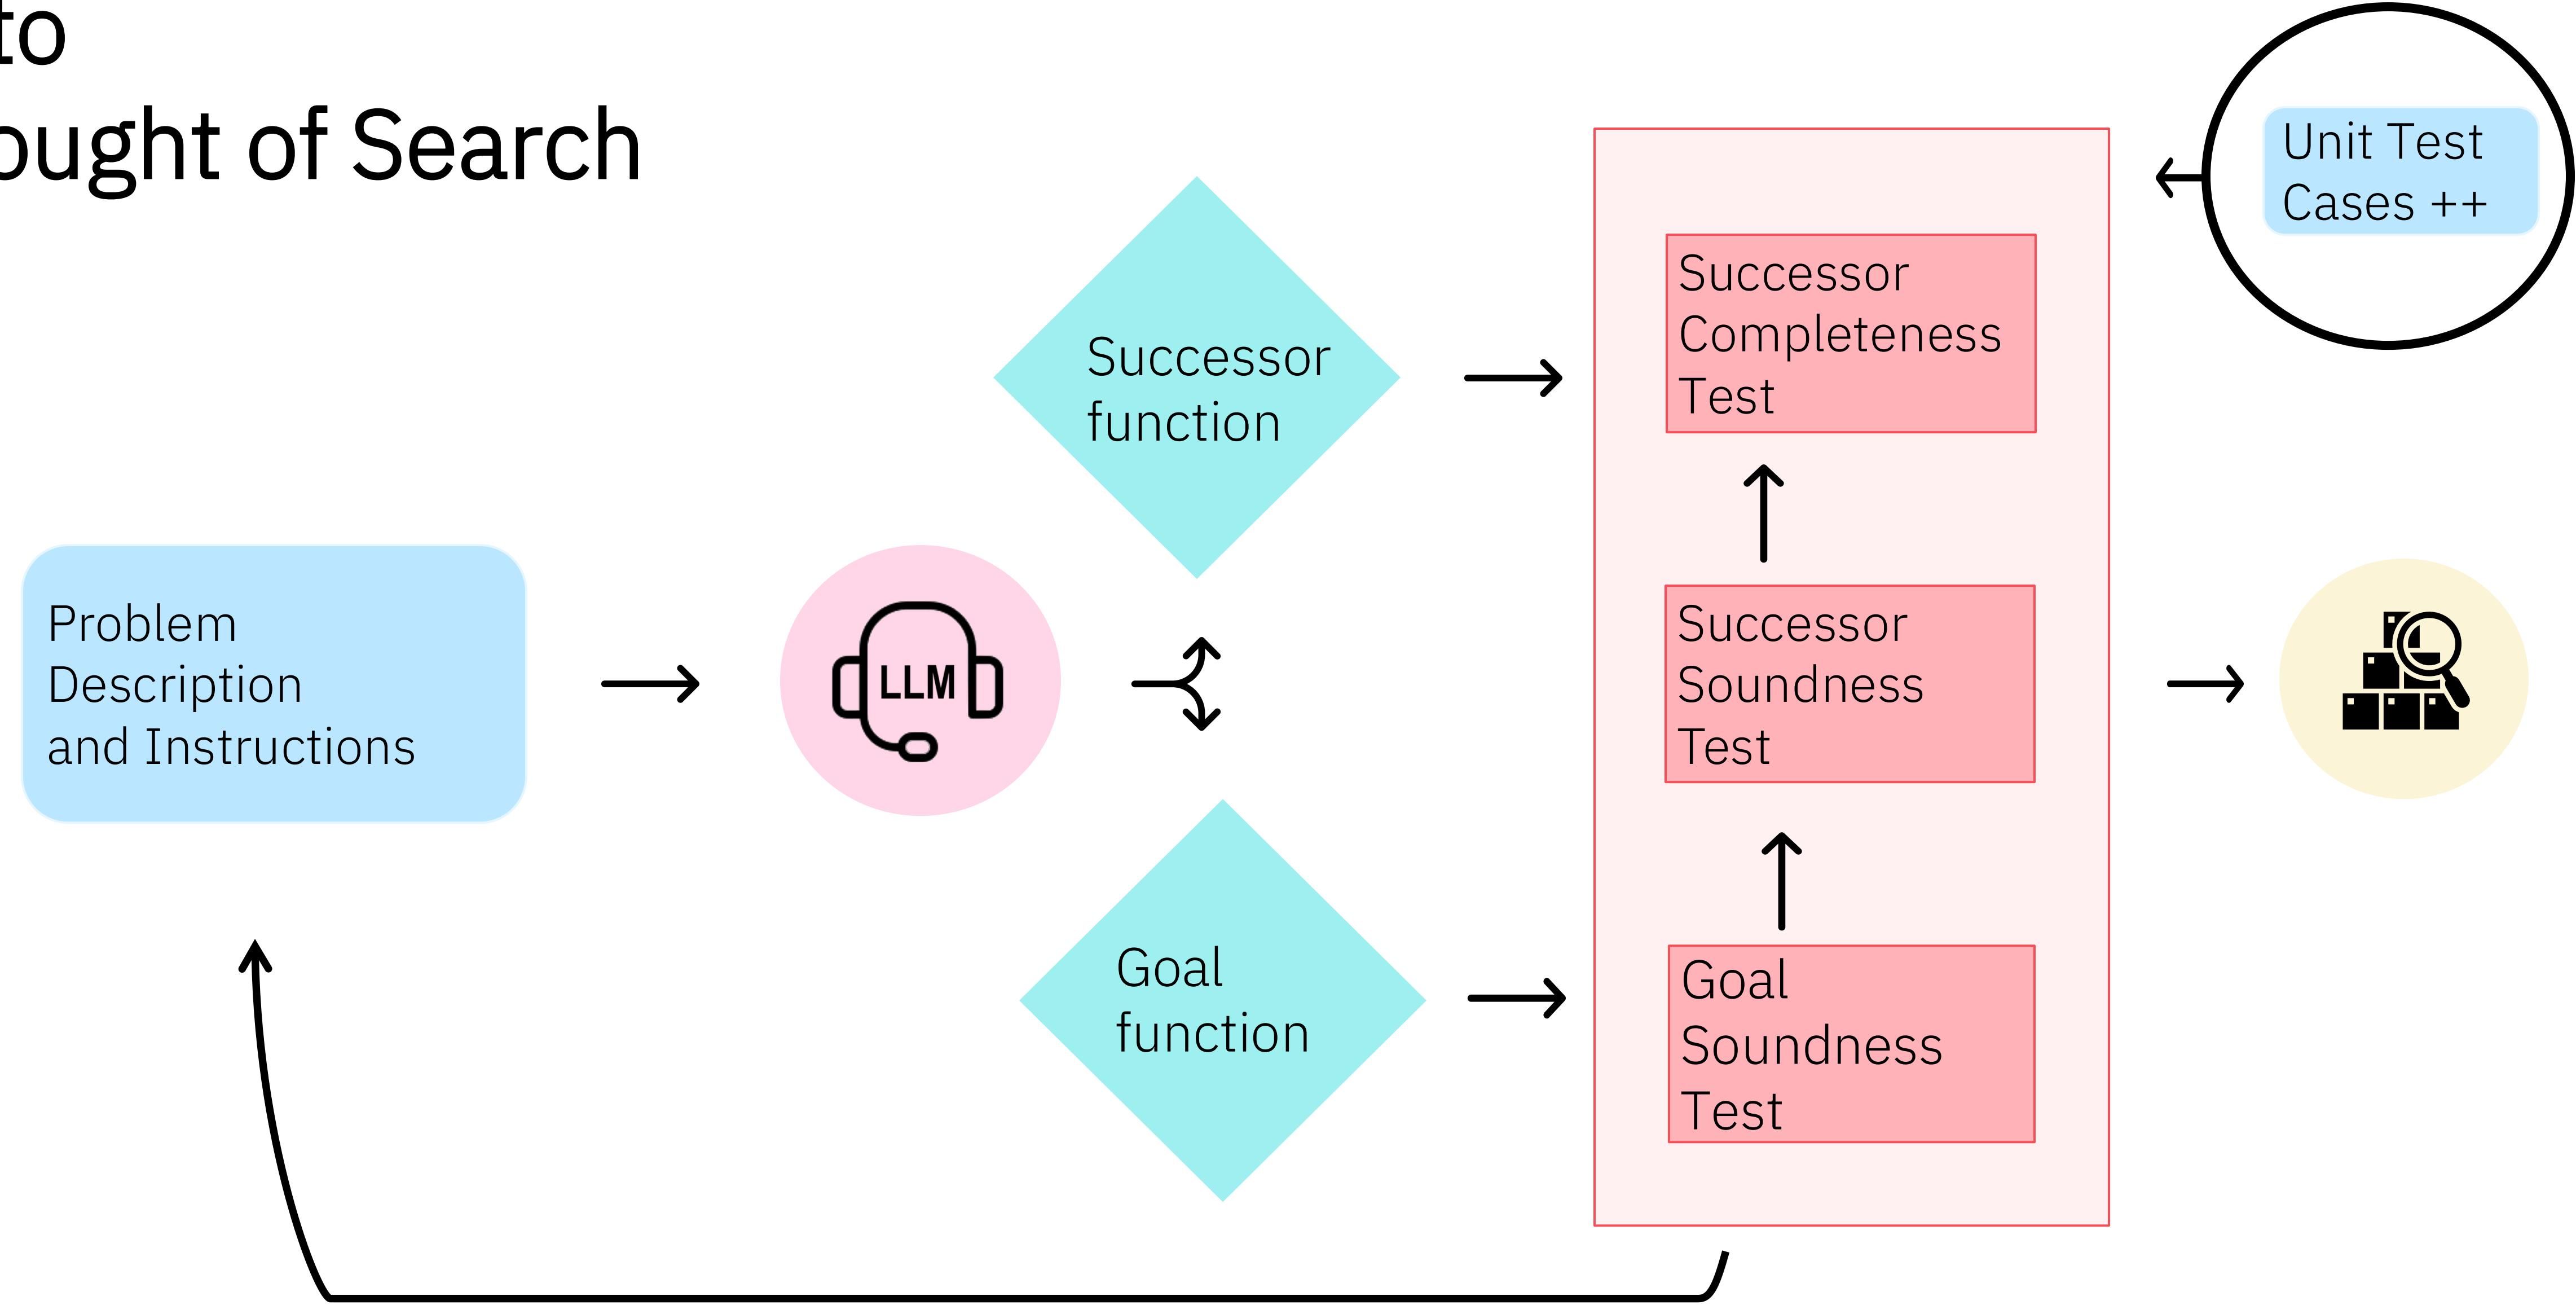
Table 1: Comparison of various approaches on evaluated tasks. The number of tasks is D , the length bound of the search/rollout/plan is L , number of rollouts is T , beam size is m , and branching bound is b . The summed number of states over the D tasks and the projected number of LLM evaluations are given per approach and problem.

100% instances solved

Thought of Search



Auto Thought of Search



Cao et al 2025

Unit Test

Goal states:

[24]

Non-goal states:

[], [3], [24, 1], [1, 6, 4], [1, 1, 4, 6]

Successors:

[[1, 1, 1, 8], [[0.125, 1, 1], [1, 1, 9], [1, 1, 8], [0, 1, 8], [1, 2, 8], [1, 1, 7], [-7, 1, 1]]]
[[6, 6, 6, 6], [[1.0, 6, 6], [6, 6, 12], [0, 6, 6], [6, 6, 36]]]]

Partial
soundness
test:

```
def validate_transition_complex(self, s, t):
    if len(s) - len(t) != 1:
        feedback = search pprint("Invalid transformation: length mismatch – the length of a successor must be
one less than the parent.")
        feedback += search pprint("Let's think step by step. First think through in words why the successor
function produced a successor that had a length that was not exactly one less than the parent. Then
provide the complete Python code for the revised successor function that ensures the length of a
successor is exactly one less than the parent.")
        feedback += search pprint("Remember how you fixed the previous mistakes, if any. Keep the same
function signature.")
        return False, feedback
    return True, ""
```

Feedback

Listing 1: 24 Game example feedback.

The goal test function failed on the following input state [24, 1], incorrectly reporting it as a goal state. First think step by step what it means for a state to be a goal state in this domain. Then think through in words why the goal test function incorrectly reported input state: [24, 1] as a goal state. Now, revise the goal test function and ensure it returns false for the input state. Remember how you fixed the previous mistakes, if any. Keep the same function signature.

Invalid transformation: length mismatch - the length of a successor must be one less than the parent. Let's think step by step. First think through in words why the successor function produced a successor that had a length that was not exactly one less than the parent. Then provide the complete Python code for the revised successor function that ensures the length of a successor is exactly one less than the parent. Remember how you fixed the previous mistakes, if any. Keep the same function signature.

Input state: [1, 1, 4, 6] Example wrong successor state: [6, 5]

Successor function when run on the state [1, 1, 4, 6] failed to produce all successors. Missing successors are: [[1, 4, 7], [-5, 1, 4], [1, 1, 2], [1, 5, 6], [0.25, 1, 6], [-3, 1, 6], [0.1666666666666666, 1, 4], [1, 3, 6], [1, 4, 5], [1, 1, 1.5]] First think step by step why the successor function failed to produce all successors of the state. Then, fix the successor function. Remember how you fixed the previous mistakes, if any. Keep the same function signature.

Experiments

	24 Game	PrOntoQA	Sokoban	Crossword	BlocksWorld
AutoToS	GPT-4o-mini	8.8	4.8	6.4	9.6
	GPT-4o	3.4	2.6	2.2	5.8
	Llama3.1-405b	3.4	2.0	2.6	4.0
	Llama3.1-70b	7.4	2.0	8.2	6.2
	DeepSeek-CoderV2	4.4	2.0	2.8	6.6
ToS	GPT-4	2.2	2.6	NA	4.2
				3.8	3.8

Table 1: The average number of calls to the language model per domain.

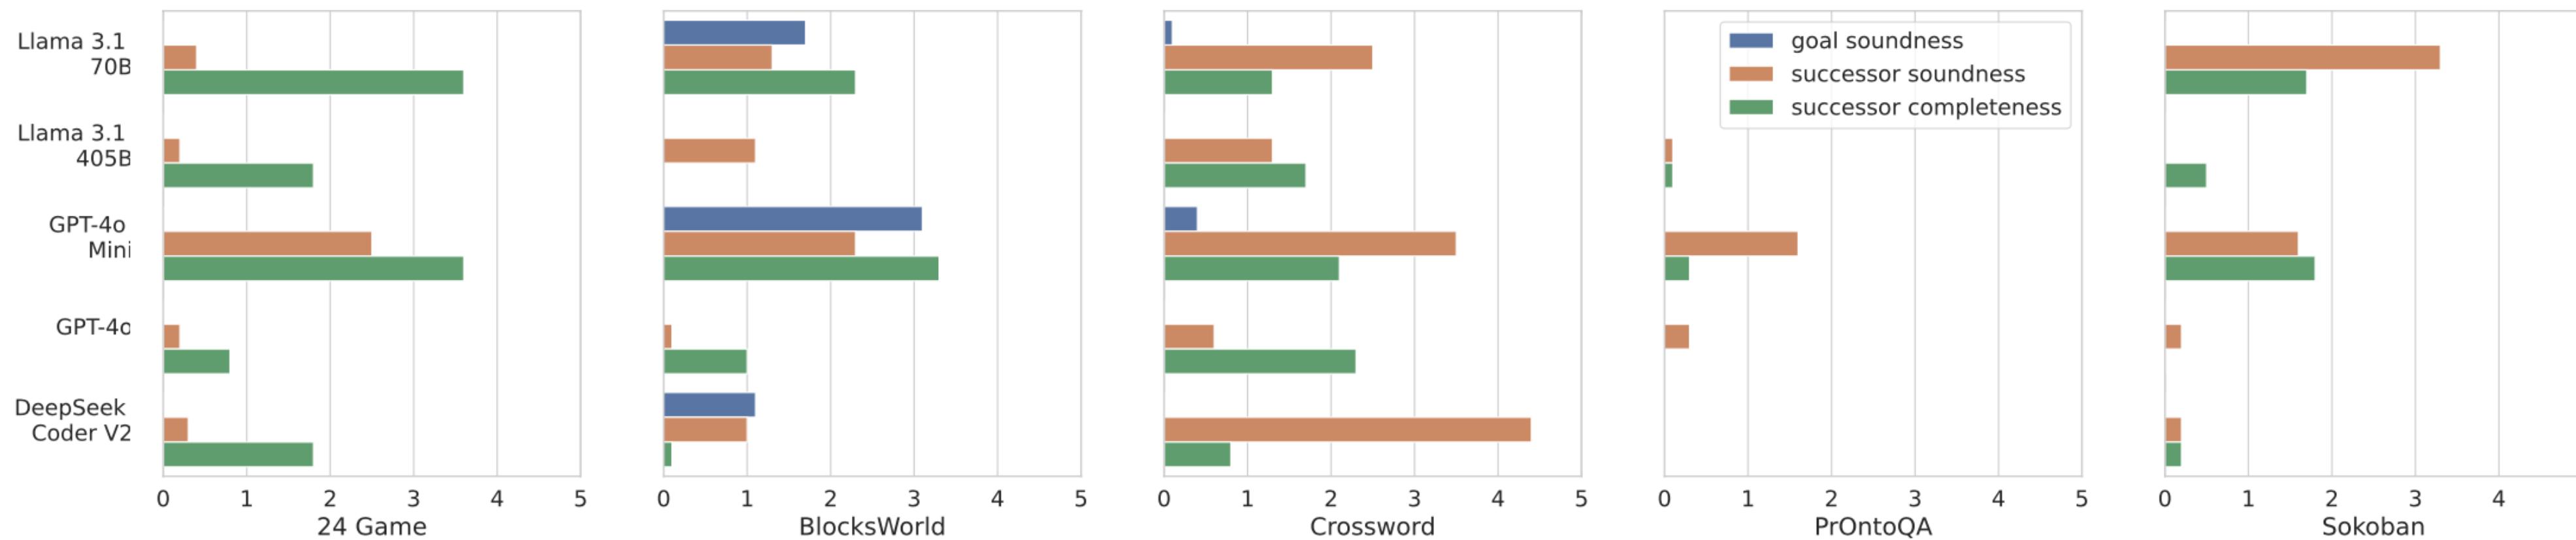
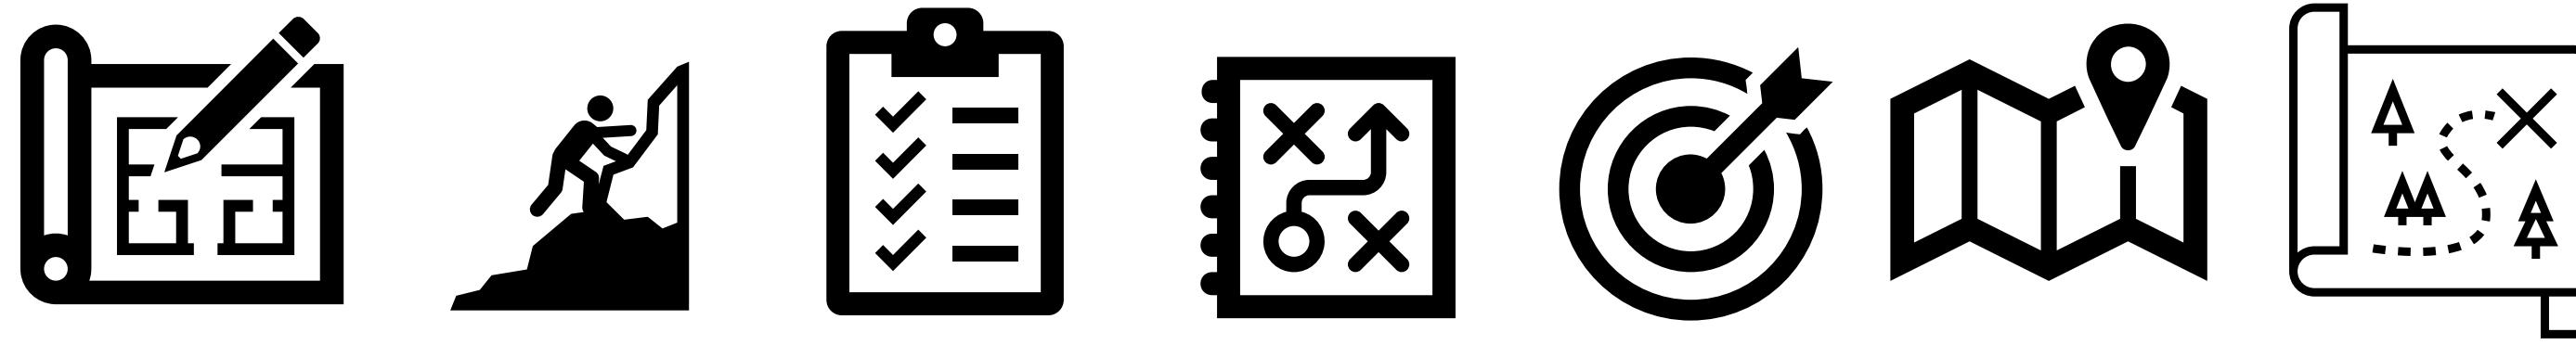


Figure 3: Average number of feedback calls for goal soundness, successor soundness, and successor completeness.

Outline

- LLMs as ~~3D Printers~~ Mold Makers
- LLMs as MoldMakers for AI Planning
 - LLM for search code generation
 - **LLM for NL2PDDL**
- Benchmarking and Evaluations

What is a Planning Problem?



Given the following description

- an initial state
- a goal (or objective)
- a set of actions that transform the state

devise a sequence of action that achieves the goal

What is a Planning Problem?

Formally,

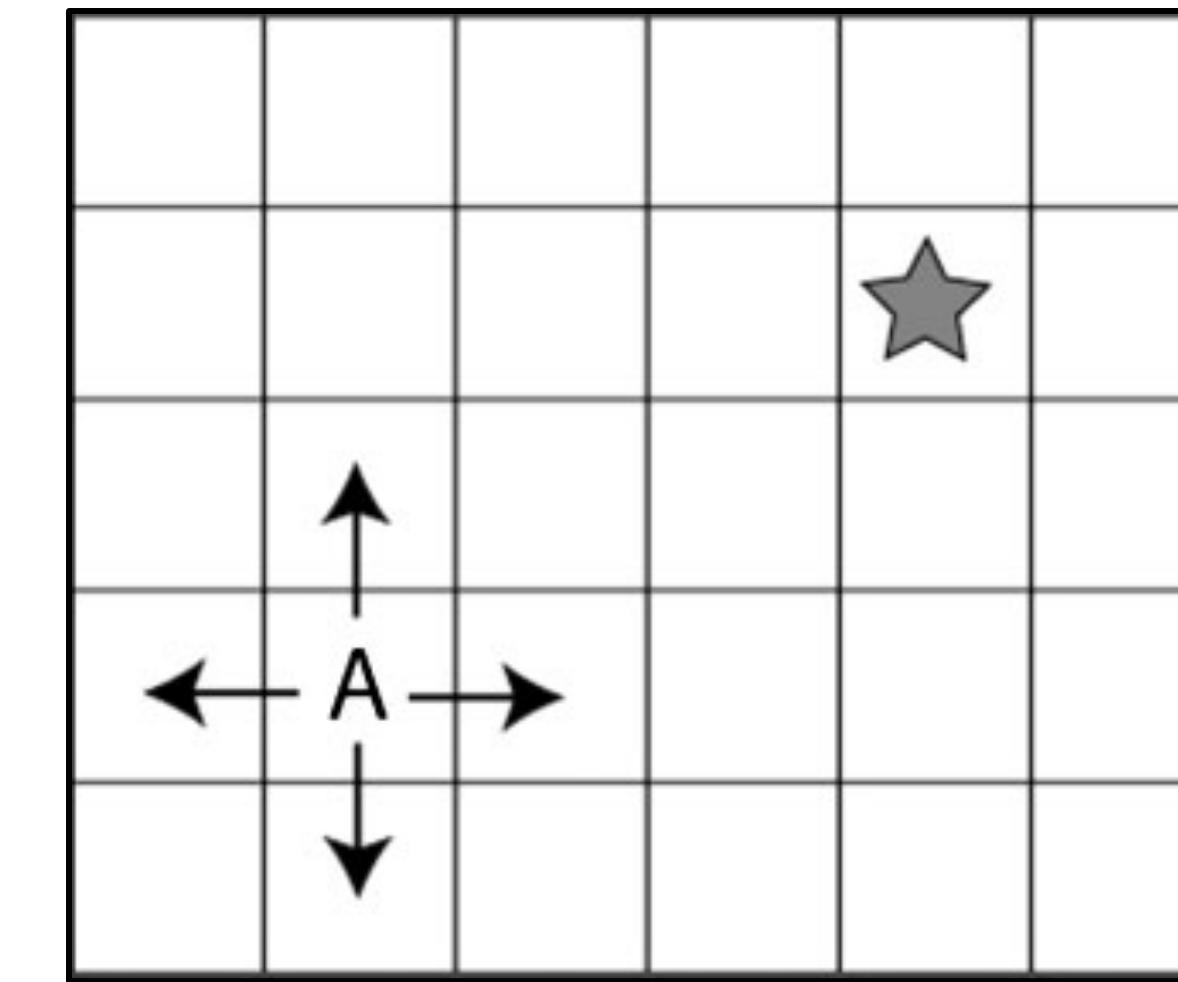
A classical planning problem includes

- a finite and discrete state space S
- a known initial state $s_0 \in S$
- a set $S_G \subseteq S$ of goal states
- a set of actions A
- a deterministic transition function $s' = f(s, a)$

PDDL – Planning Domain Definition Language

```
( define (domain grid)
  (:requirements :strips)
  (:types place - object)
  (:predicates (adj ?x ?y)
    (at-robot ?x))

  ( :action move
    :parameters (?from – place ?to – place)
    :precondition (and (at-robot ?from)
      (adj ?from ?to ))
    :effect (and (at-robot ?to)
      (not (at-robot ?from)))
  )
)
```

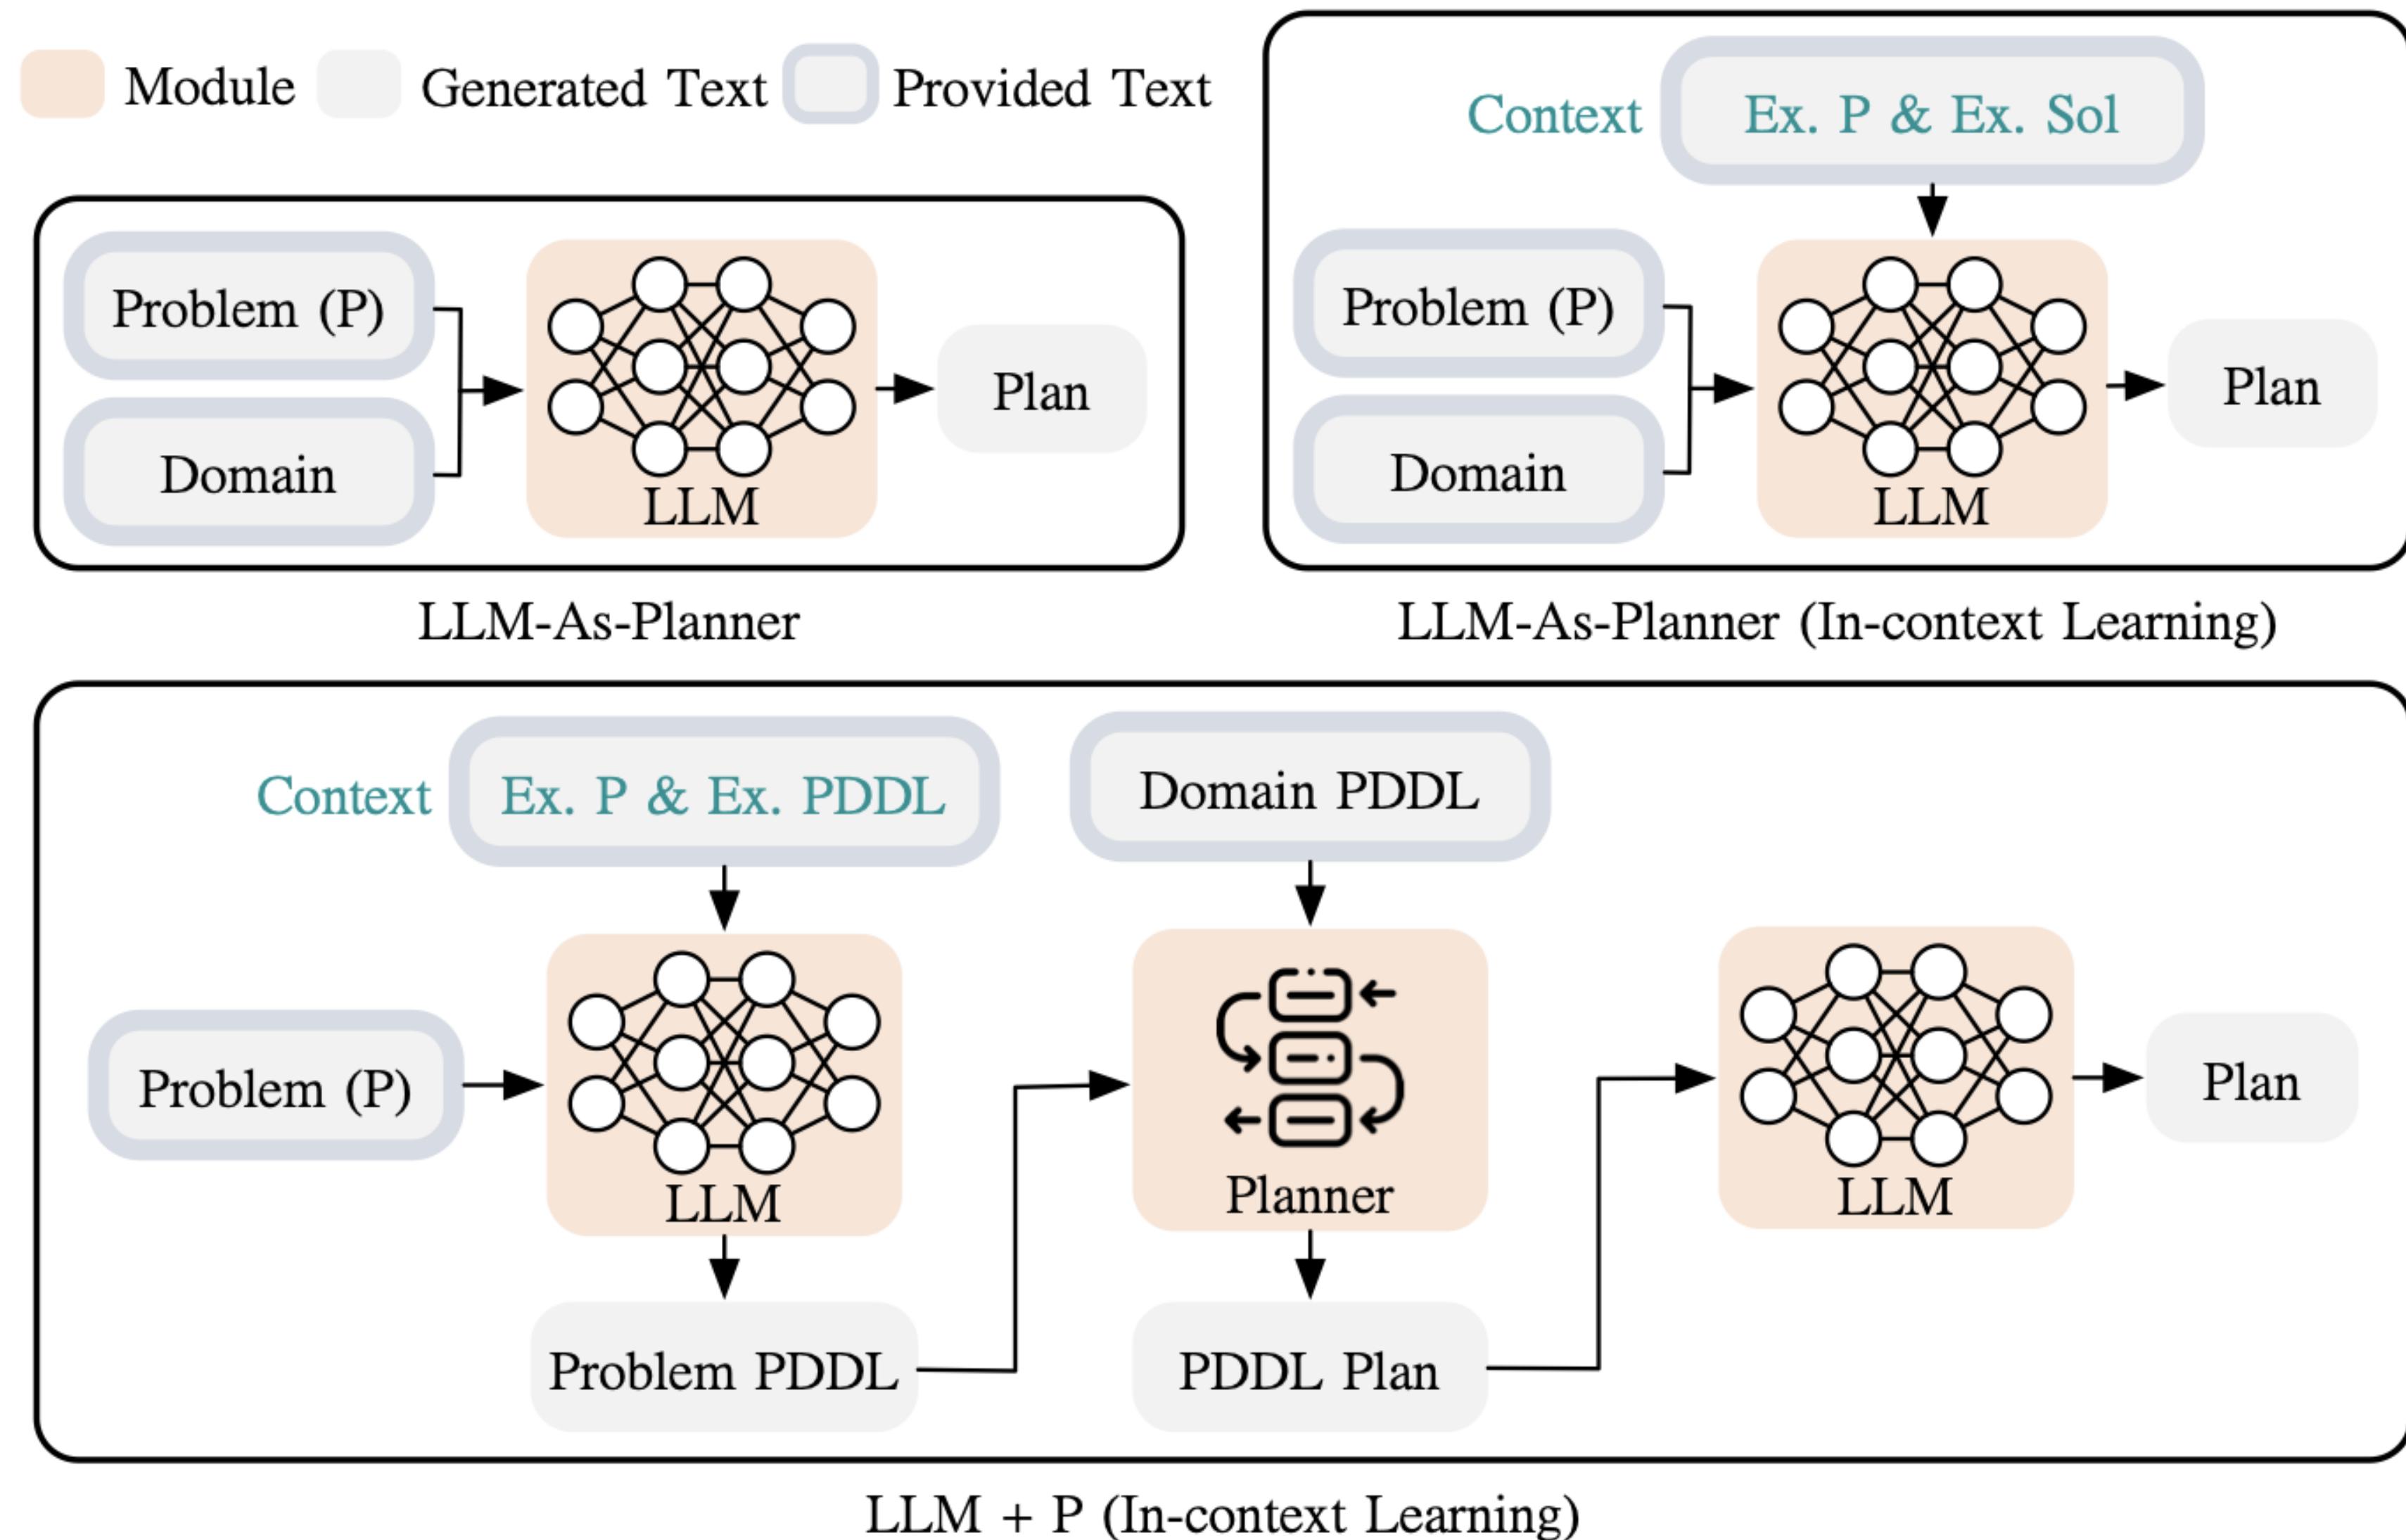


```
( define (problem small)
  (:domain grid)
  (:objects node0-0 node0-1 ...)
  (:init (at-robot node1-1)
    (adj node0-0 node0-1) ...)
  (:goal (at-robot node3-4))
)
```

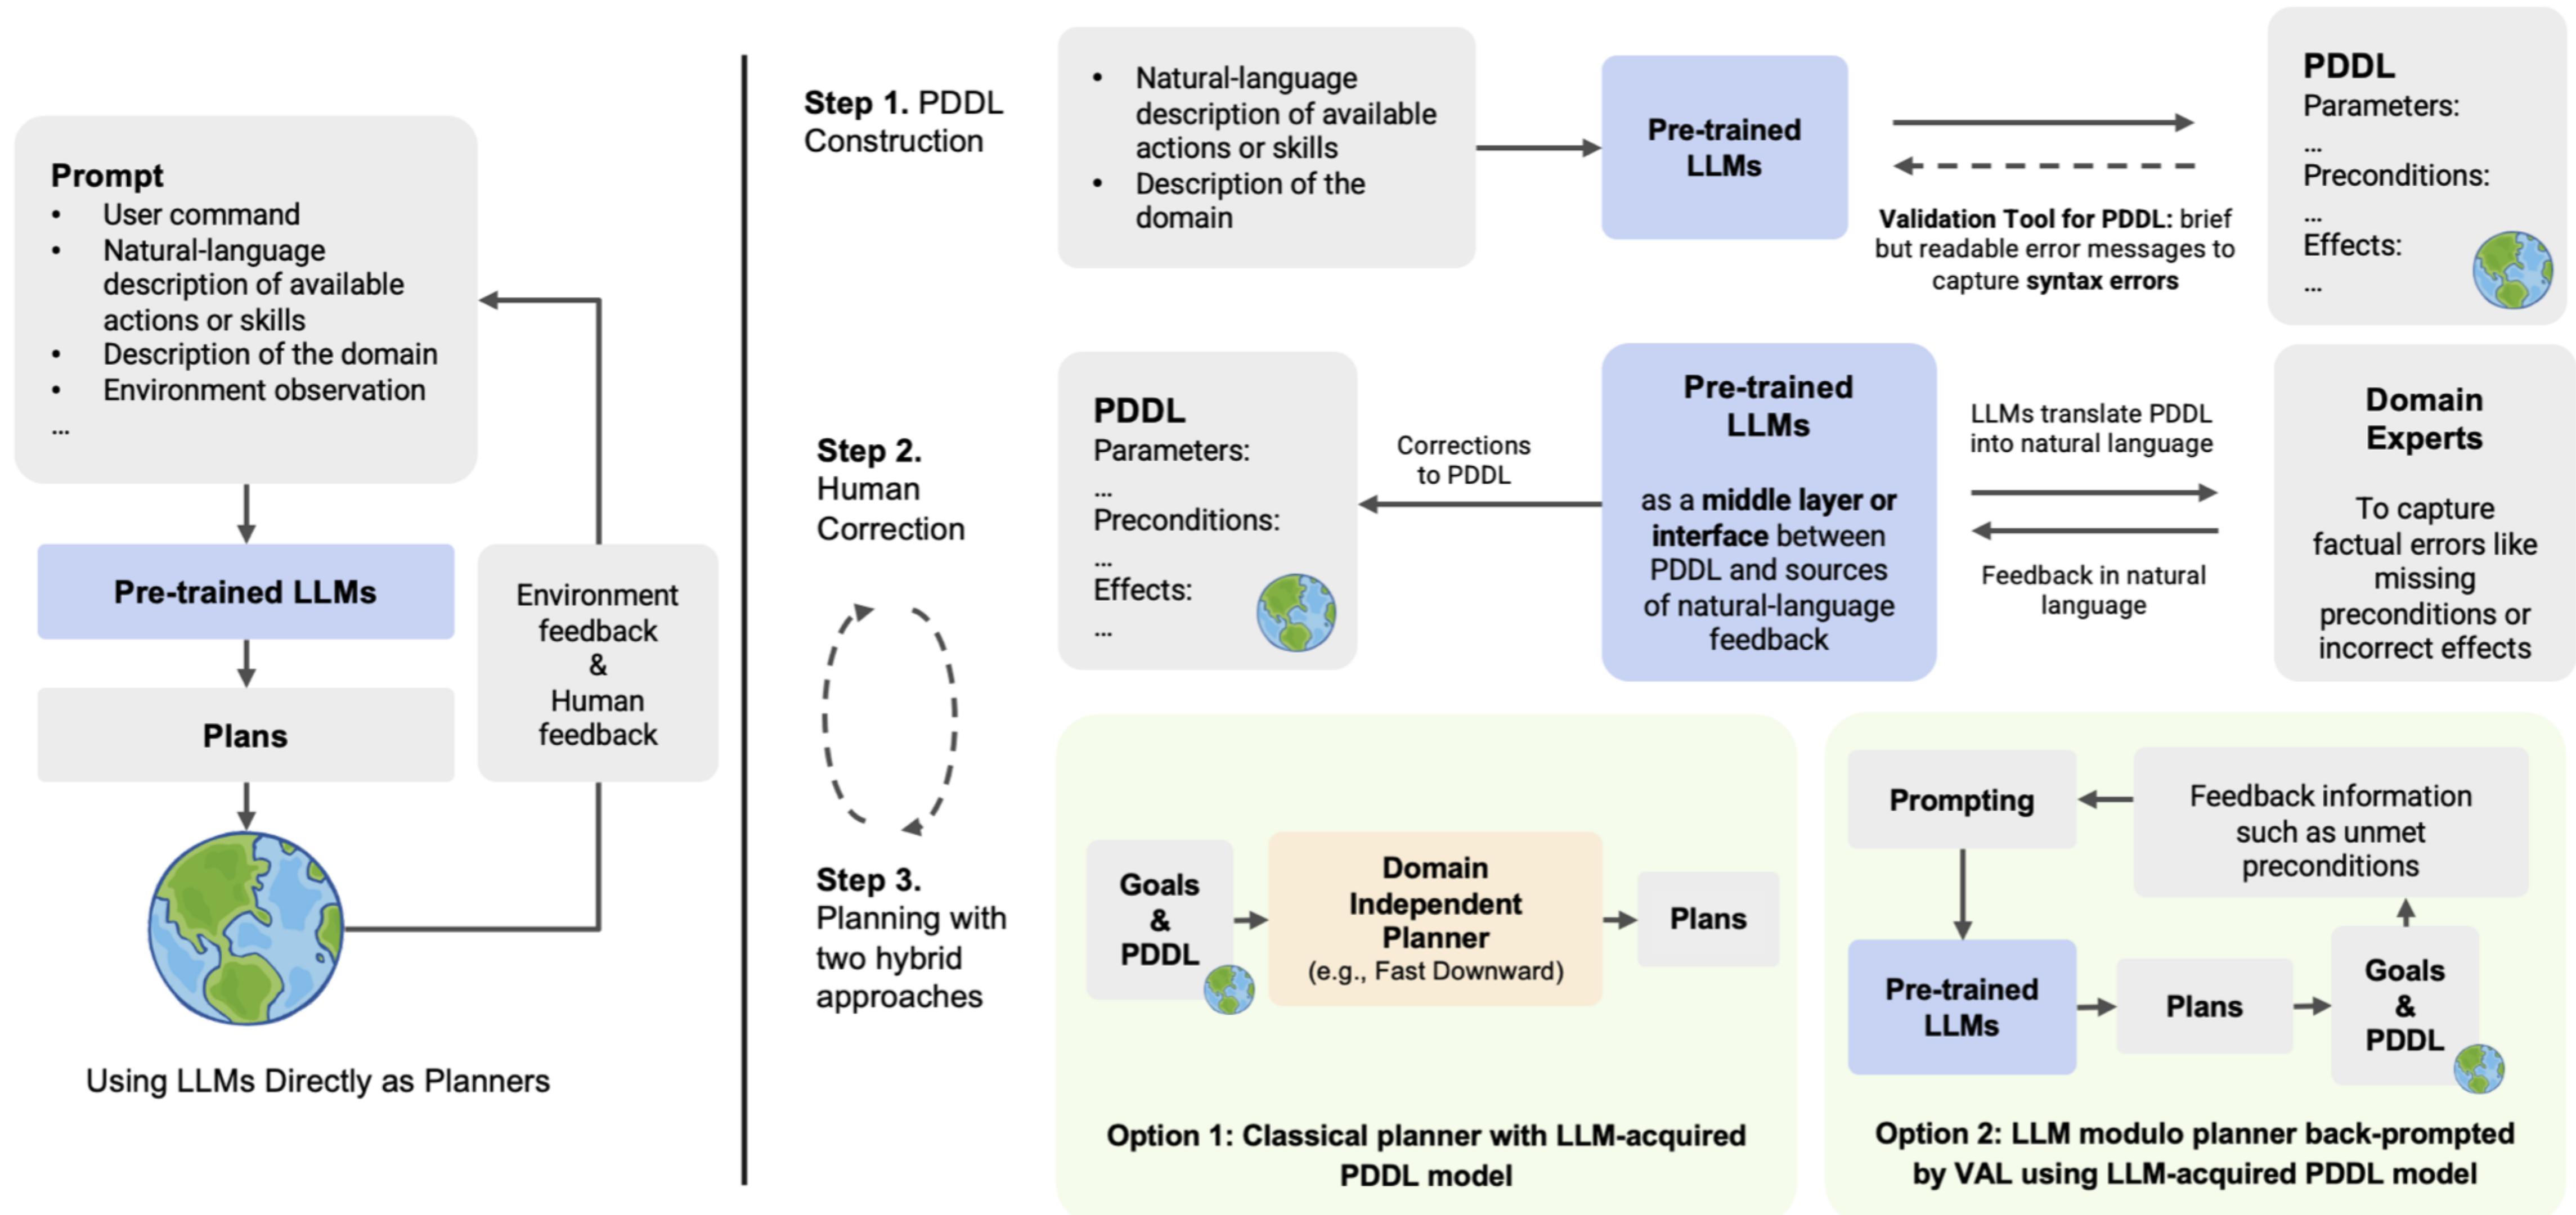
Why PDDL?

- **Formal Analysis**
- **Standardization:**
Common language for International Planning Competition (IPC)
- **Clarity and Precision:**
important for soundness
- **Tool Compatibility:**
Allows domain-independent planning systems

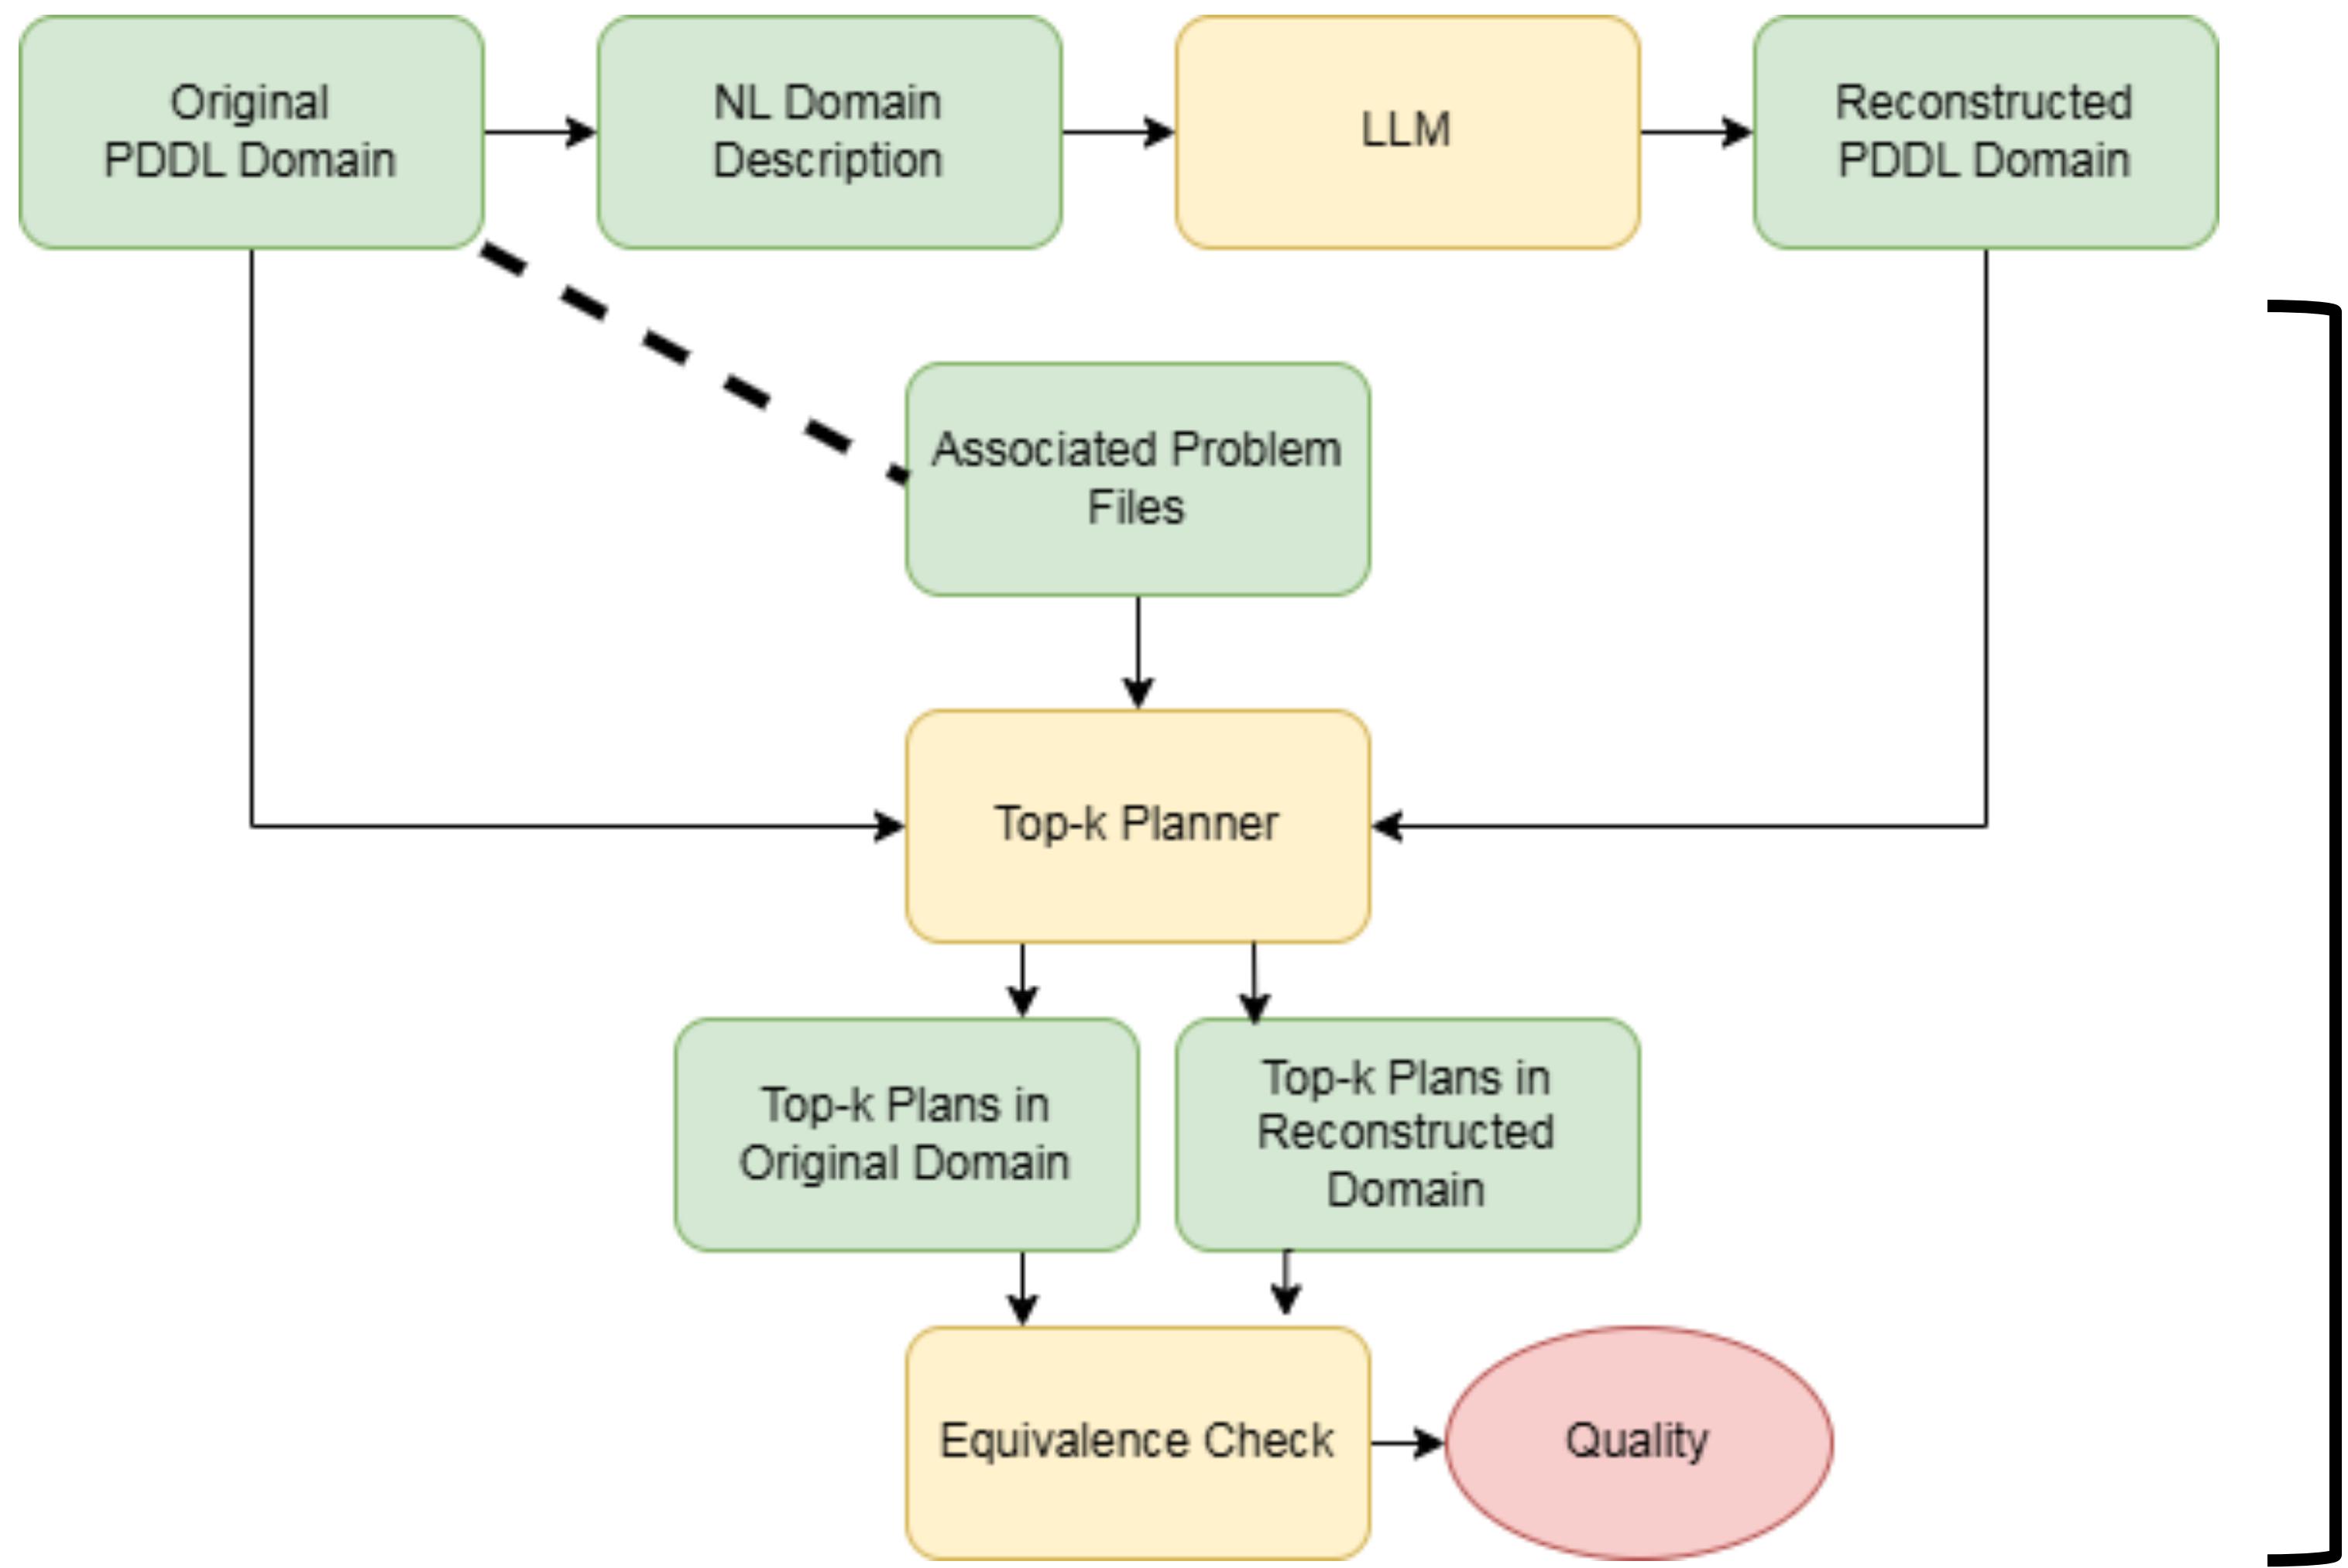
LLM+P



LLM-DM



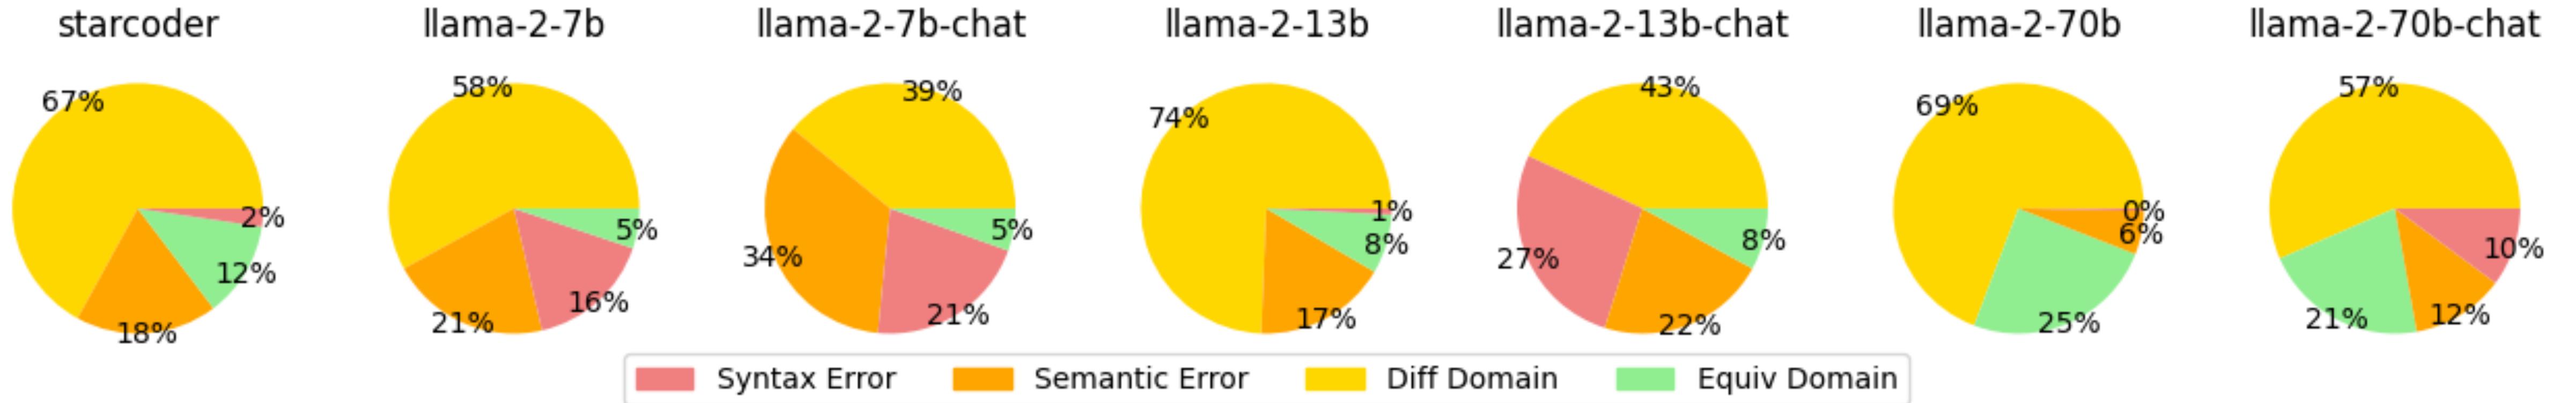
NL2PDDL



NL2PDDL

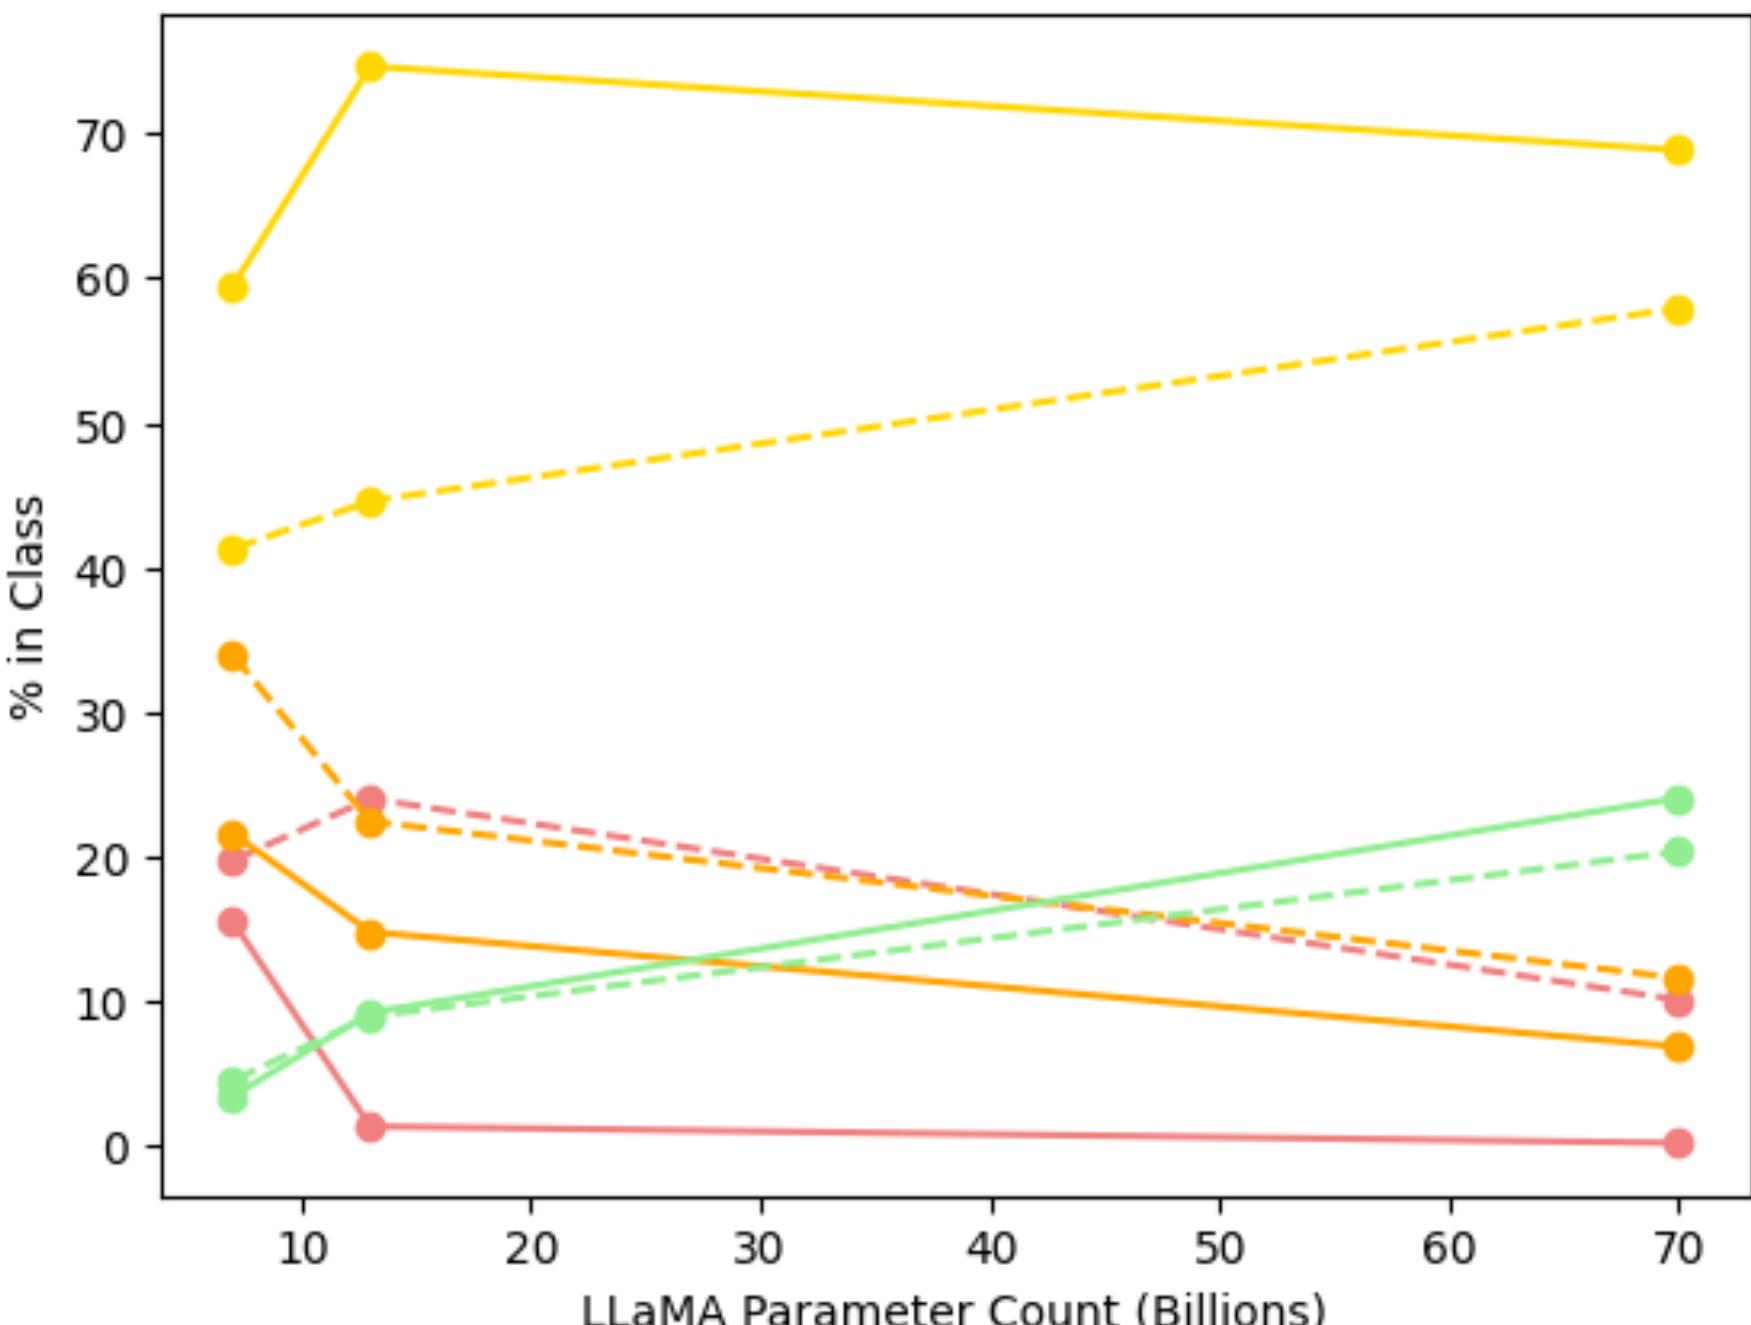
Error/Result Classes

- **Syntax Error:** The model produced syntactically invalid PDDL.
- **Semantic Error:** The model produced syntactically valid PDDL, but the PDDL doesn't integrate with the intended problems due to type mismatches, wrong number of parameters to predicates/actions, etc.
- **Different Domain:** The model produced syntactically valid PDDL that integrates with the original domain, but the underlying domains are different because they produce different plans for the same problem.
- **(Heuristically) Equivalent Domain:** The model produced syntactically valid PDDL that integrates with the desired domain. Plans from the original domain can be applied in the new domain and vice versa.



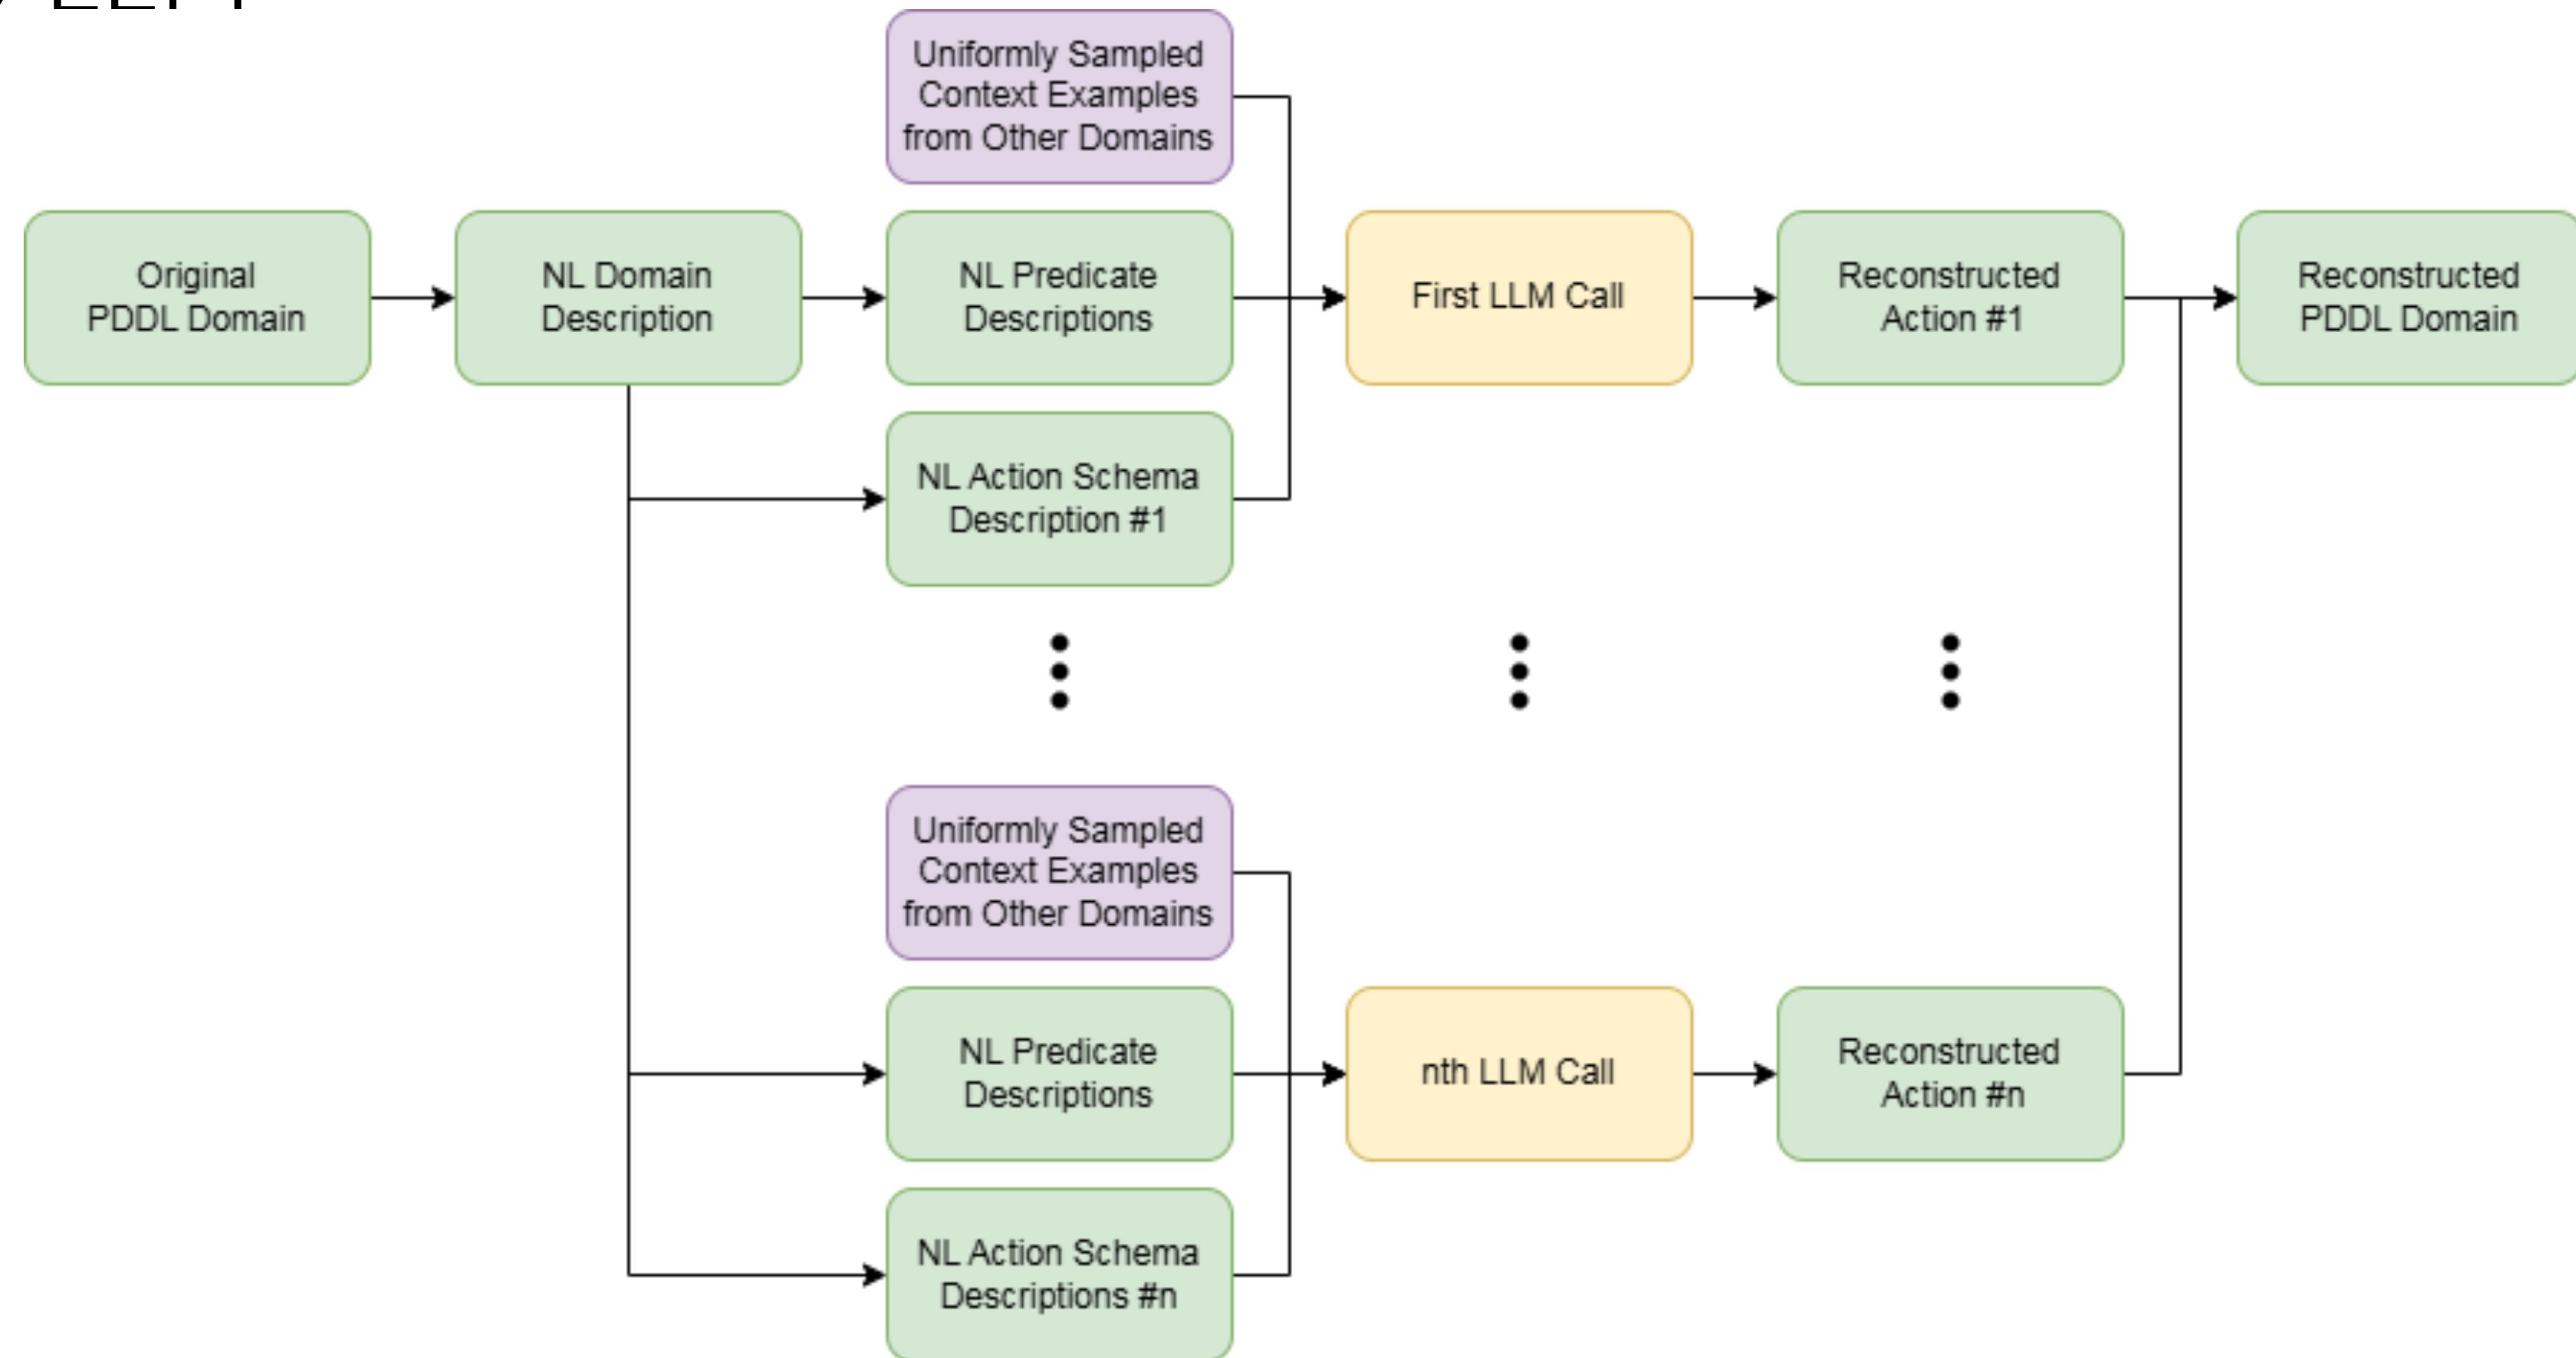
- ▶ Blocksworld : 5 preds 4 actions
- ▶ Depot : 6 preds 5 actions
- ▶ Forest : 5 preds 2 actions
- ▶ Logistics : 3 preds 6 actions
- ▶ Miconic : 6 preds 4 actions
- ▶ TrapNewspapers : 7 preds 3 actions
- ▶ Heavy-pack* : 5 preds 2 actions
- ▶ Trackbuilding* : 4 preds 3 actions

* : Domain is not present in training data.



NL2PDDL

Call to LLM



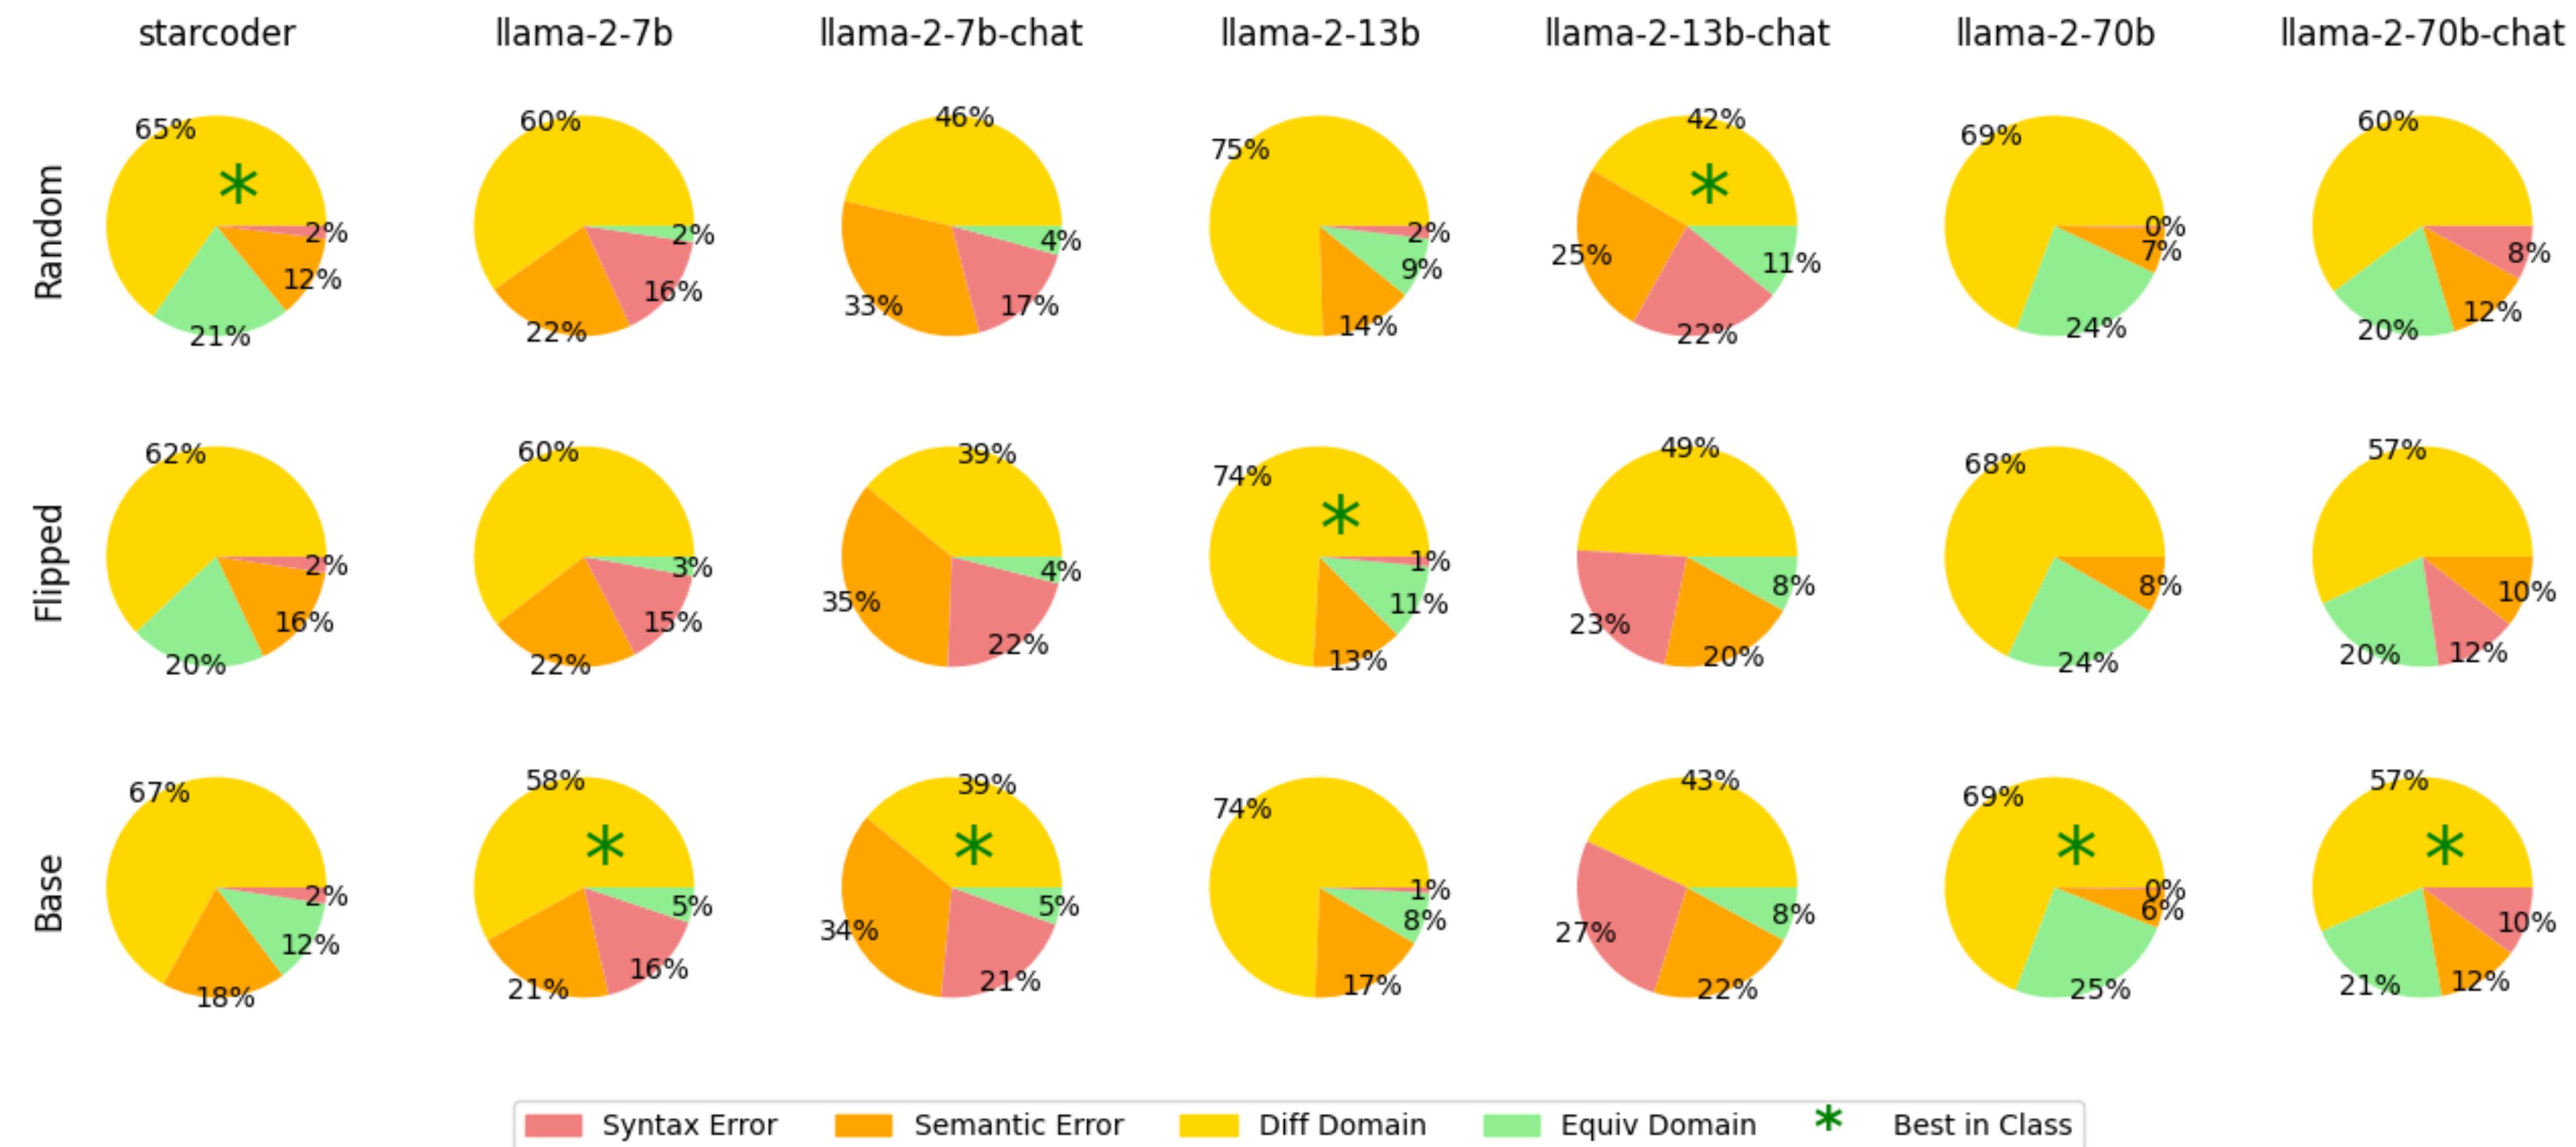
NL2PDDL

Prompt Variants

- ▶ **Base:** “*The action ‘unstack’ will have a hand unstack a block x from a block y.*”
- ▶ **Flipped:** “*The action ‘unstack’ will have a hand unstack a block x from a block y, if the block x is clear, x is on y, and the hand is empty.*”
- ▶ **Random:** “*The action ‘unstack’ will have a hand unstack a block x from a block y, if the hand is empty and x is on y. After the action, y should be clear.*”

NL2PDDL

Prompt Variants



NL2PDDL

Sample Prompt

Instruction

Given a description of an action in some domain, convert it to Planning Domain Definition Language (PDDL) action. You may only use the allowed predicates provided for each action.

Context

Allowed Predicates:

(agent-at ?loc - location) : the agent is at the location loc.

{Remaining Predicates for context example 1}

Input:

The action, "build-track" will build track at a location.

PDDL Action:

```
(:action build-track
  :parameters (?loc - location)
  :precondition (agent-at ?loc)
  :effect (and (has-track ?loc))
)
```

{2 Other Context Examples from different domains}

Task

Allowed Predicates:

(at ?x - locatable ?y - place) : the locatable x is at some place y.

{Remaining Predicates for Task}

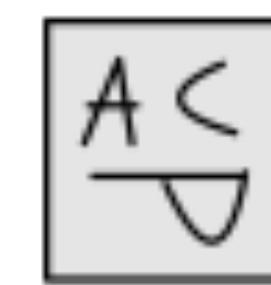
Input:

The action, "Unload" will use a hoist to unload a crate from a truck at a place.

PDDL Action:

Benchmarking and Evaluation





ACPBENCH

Next Action
Landmarks
Justification
Action Reachability
Action Applicability
Reachability
Progression
Validation

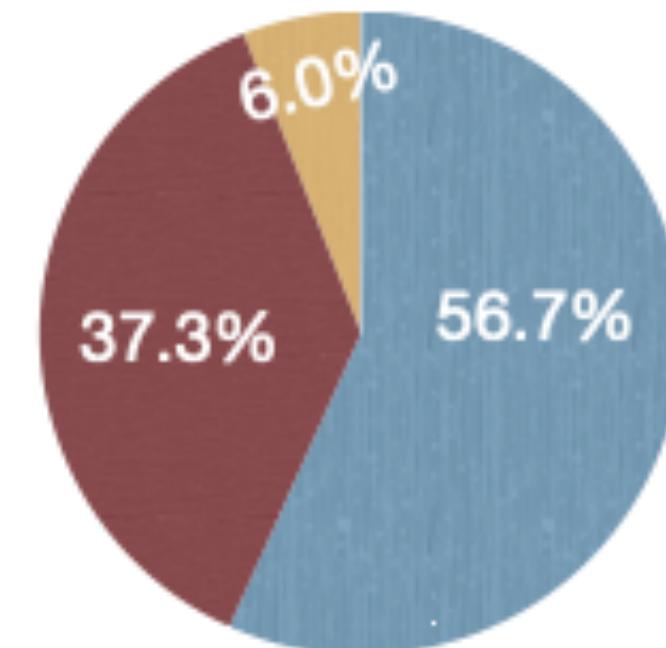
8 reasoning tasks

FerryLogistics
SwapGripper
AlfWorld
BlocksWorld
Goldminer
GridSatellite

13 Domains*

* can add more domains in a day

Action Applicability



```
{  
  "id": 2606453784296512791,  
  "group": "applicable_actions_bool",  
  "context": "This is a ferry domain, where the task is to transport cars from their start to their goal locations, using a ferry. Each location is accessible by ferry from each other location. The cars can be debarked or boarded, and the ferry can carry only one car at a time. There are 3 locations and 10 cars, numbered consecutively. Currently, the ferry is at l1, with the car c2 on board. The cars are at locations as follows: c6, c3, and c0 are at l2; c4, c9, and c7 are at l0; c1, c8, and c5 are at l1.",  
  "question": "Is the following action applicable in this state: debark the car c2 from the ferry to location l1?"  
}
```

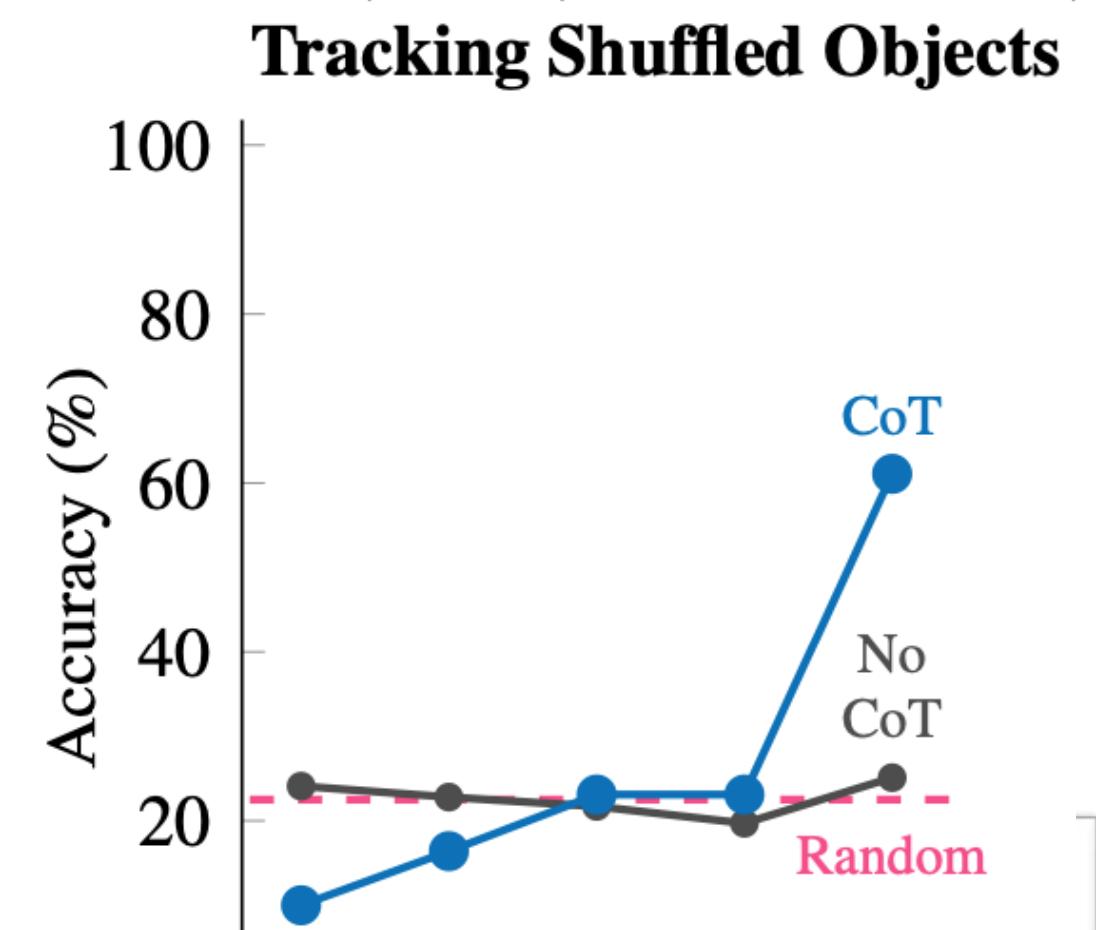
Xie et al. TravelPlanner

Progression

```
{
  "id": 2142145895175892935,
  "group": "progression_bool",
  "context": "This is a ferry domain start to their goal locations, using each other location. The cars carry only one car at a time. They consecutively. Currently, the ferris at locations as follows: c9, c4, are at l1; c3 and c5 are at l2.",
  "question": "Will the fact \"The fact \"embark the car c0 at location l1
}
```

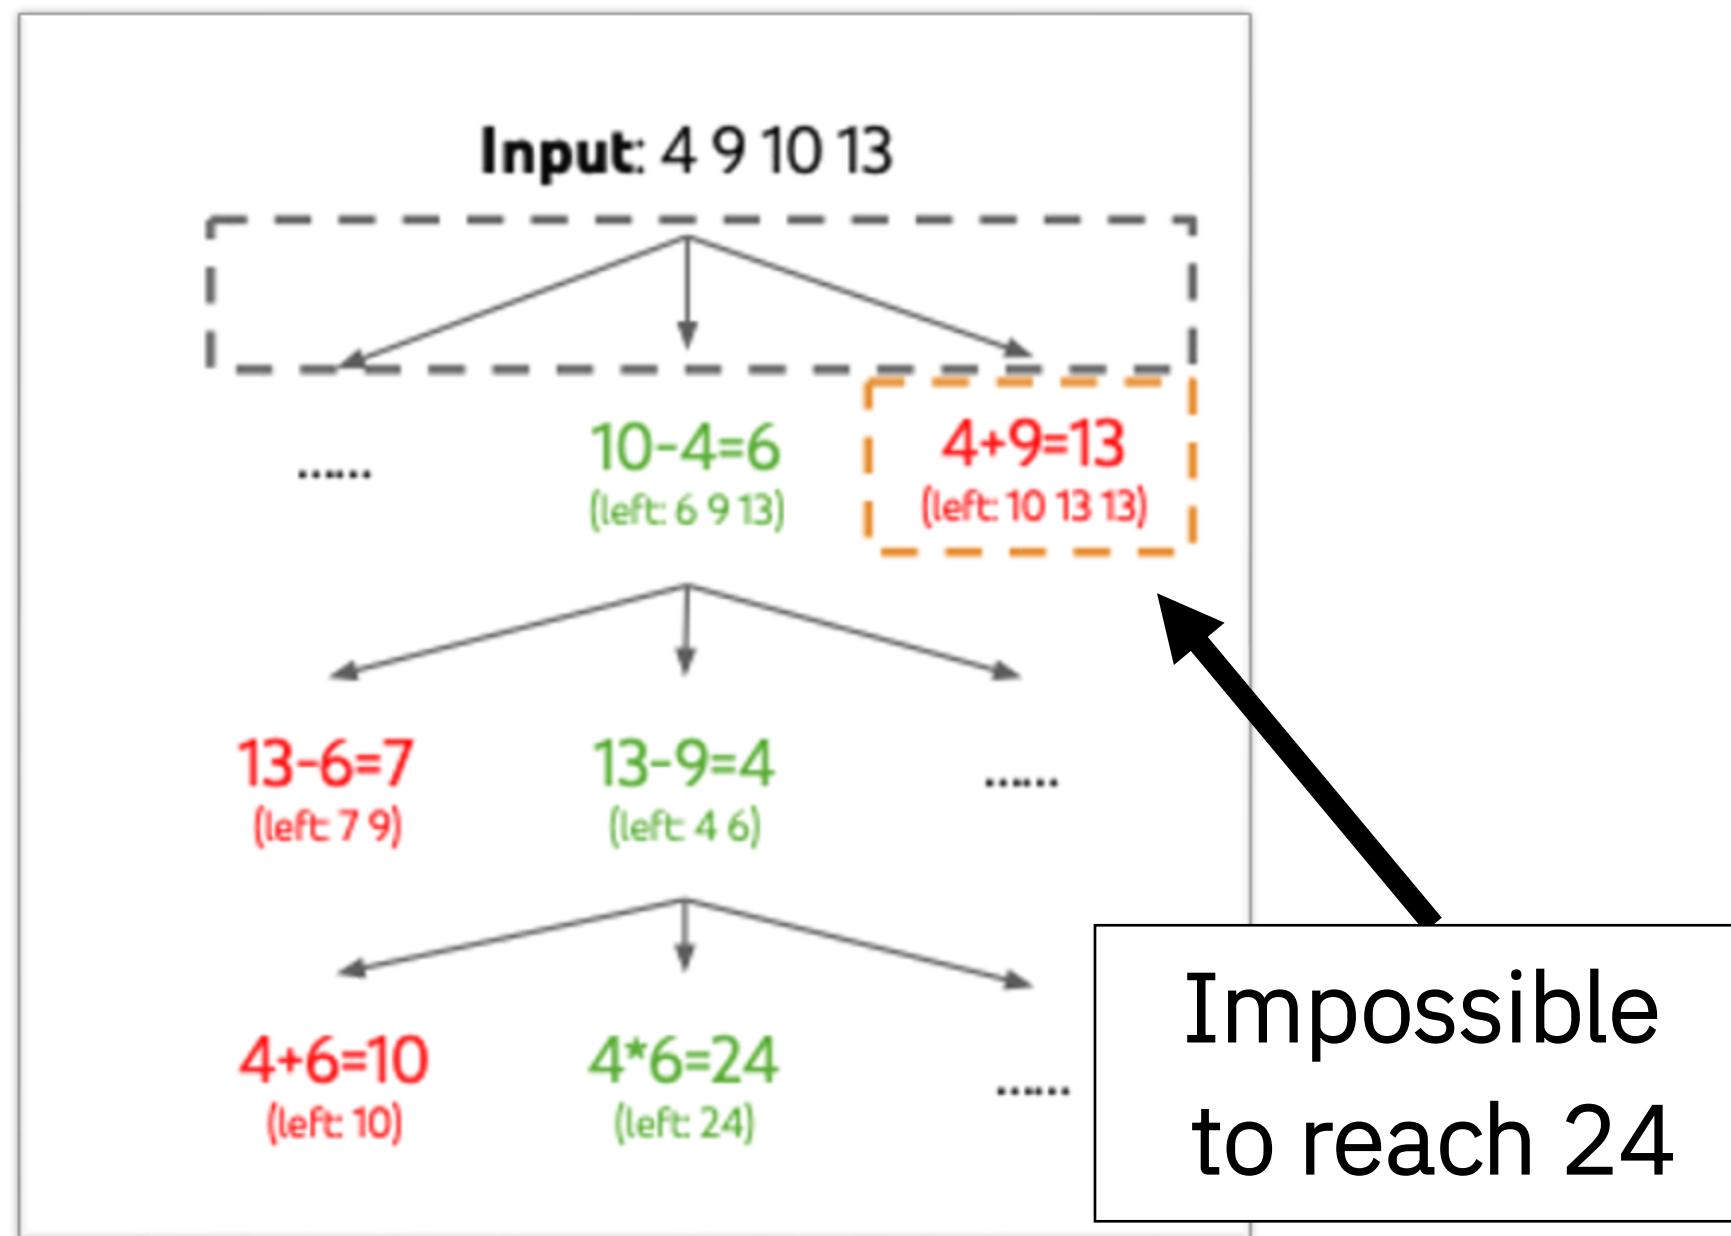
Suzgun et al.

BigBenchHard



Break down the outcomes of performing the action \"sail from location l4 to location l0\" into two lists, positive effects and negative effects. Positive effects are the propositions that are false in the current state but will become true after performing the action. Negative effects are the propositions that are true in the current state and will become false after performing the action. Provide only the two lists with the ground propositions.

Reachability



Yao et al. Tree of Thoughts

```
id": -3047022720995966296,  
group": "reachable_atom_bool",  
context": "This is a ferry domain, where the task is to transport cars from the start to their goal locations, using a ferry. Each location is accessible by ferry from each other location. The cars can be debarked or boarded, and the ferry can carry only one car at a time. There are 5 locations and 3 cars, numbered consecutively. Currently, the ferry is at l0, with the car c1 on board. The cars are at locations as follows: c0 and c2 are at l3.",  
question": "Is it possible to transition to a state where the following holds: Ferry is empty and The ferry is at l2 location?"
```

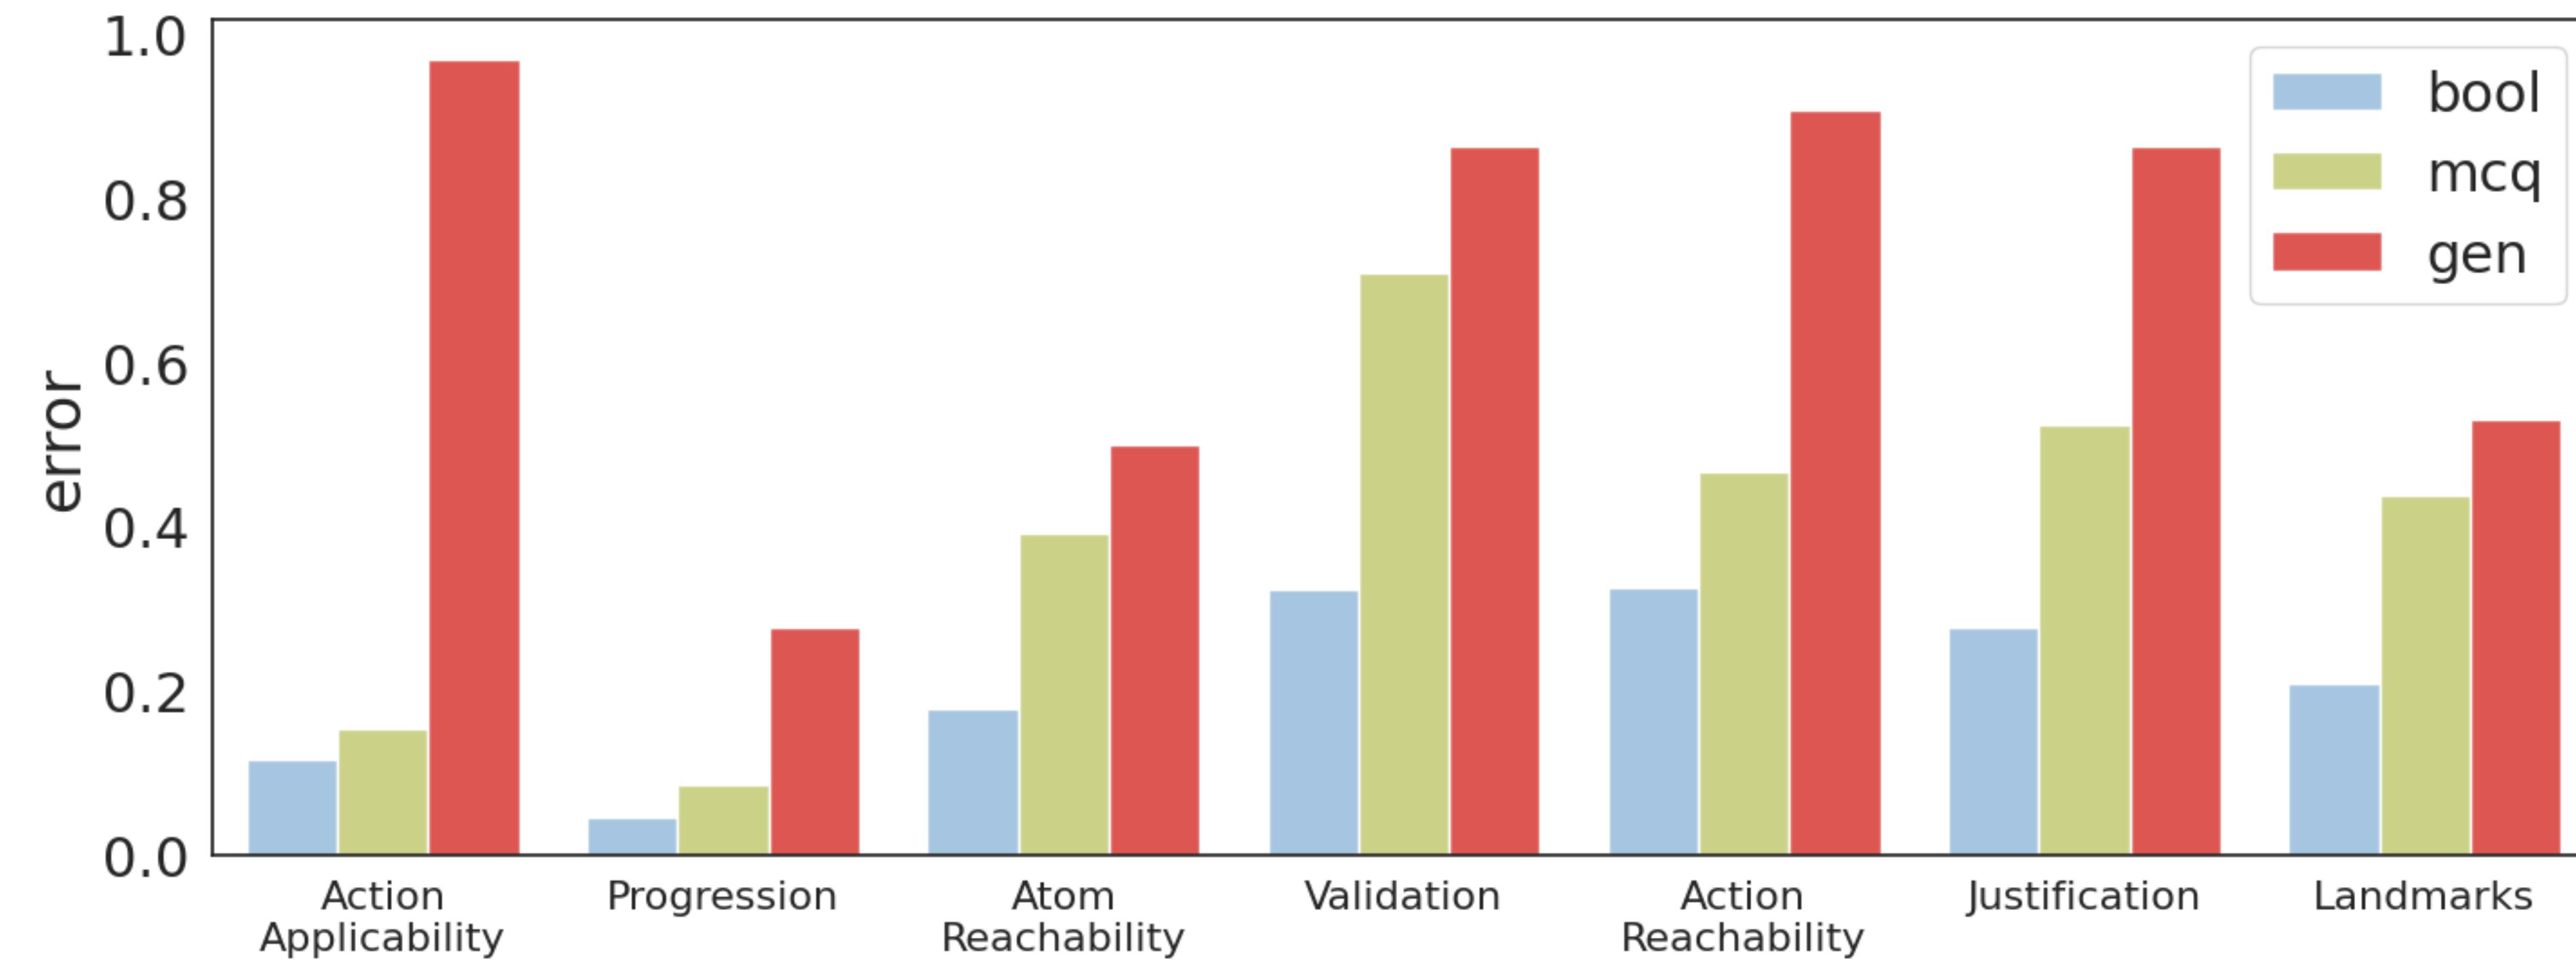
```
"choices": {  
    "text": [  
        "Car c29 is on board the ferry and The ferry is empty",  
        "Car c30 is at location c36",  
        "The ferry is at l3 location",  
        "The ferry is at c30 location and Car c2 is at location l1"  
    ],  
    "label": [  
        "A",  
        "B",  
        "C",  
        "D"  
    ]  
},  
"query": "Which fact is reachable from this state?"
```

Dataset	PlanBench	AutoPlanBench	TRAC	ARB	ACP Bench	ACP Bench Hard
# Tasks	8	1	4	6	7	8
# Domains	3 (+variants)	13	1	8	13	13
NL templates	✓	✗	✓	✓	✓	✓
Evaluation	✗	✗	↔	↔, LLM	↔	✗
Question Format						
Generative	✓	✓	✗	✓	✗	✓
Boolean	✗	✗	✓	✓	✓	✗
MCQ	✗	✗	✗	✗	✓	✗
Tasks						
Applicability	✗	✗	✓	✓	✓	✓
Progression	✓	✗	✓	✓	✓	✓
Reachability	✗	✗	✗	✗	✓	✓
Action Reachability	✗	✗	✗	✗	✓	✓
Validation	✓	✗	✓	~	✓	✓
Justification	✗	✗	✗	✗	✓	✓
Landmark	✗	✗	✗	✗	✓	✓
Next Action	✗	✗	✗	✗	✗	✓

Table 4: Comparison of ACPBench-hard with existing Planning Benchmarks. Evaluations are either using string matching (↔), symbolic tools (✗), or using another LLM (LLM).

Model	Applicability		Progression		Reachability		Validation		Action Reach.		Justification		Landmark		Mean	
	Bool	MCQ	Bool	MCQ	Bool	MCQ	Bool	MCQ	Bool	MCQ	Bool	MCQ	Bool	MCQ	Bool	MCQ
Phi-3 128K	66.15	33.08	68.46	53.85	52.31	26.15	50.77	19.23	53.33	32.50	49.23	33.85	49.23	46.92	55.53	34.75
Gemma 7B	63.23	28.62	64.92	31.08	53.08	23.08	46.92	20.0	55.67	34.50	50.77	36.46	27.54	30.31	51.80	28.93
Mistral 7B	61.54	32.31	73.08	38.46	53.08	28.46	47.85	17.69	65.00	19.17	48.46	30.00	35.38	33.08	55.00	28.67
Mistral I. 7B	63.08	31.54	61.54	46.92	61.54	33.08	52.15	36.15	45.83	34.17	43.08	29.23	57.69	50.77	55.45	37.30
Granite C. 8B	59.23	32.31	70.00	34.31	52.31	24.31	44.15	17.08	57.50	25.83	46.92	34.62	37.23	35.38	53.09	29.21
Granite 3.0 8B	72.31	26.92	73.08	53.85	53.08	24.62	53.08	20.00	45.83	30.83	49.23	34.62	42.31	34.62	55.56	32.21
Granite 3.0 I. 8B	76.92	30.00	73.85	57.69	53.08	36.92	55.38	34.62	58.33	44.17	<u>70.77</u>	31.54	51.54	43.08	62.84	39.72
LLAMA-3 8B	72.92	49.23	73.08	56.00	55.23	41.08	51.54	<u>49.23</u>	<u>63.50</u>	36.67	<u>57.54</u>	32.31	56.92	43.85	61.53	44.05
LLAMA-3.1 8B	65.38	56.92	63.85	47.69	53.08	33.85	60.00	<u>37.69</u>	42.50	28.33	46.92	45.38	33.85	40.00	51.46	41.52
Mixtral 8x7B	<u>75.85</u>	<u>57.69</u>	<u>74.00</u>	<u>61.38</u>	<u>76.00</u>	40.00	65.69	34.77	52.83	<u>55.00</u>	55.38	51.38	59.54	<u>60.00</u>	65.53	<u>51.44</u>
Codestral 22B	<u>84.62</u>	<u>39.23</u>	<u>83.85</u>	<u>51.54</u>	<u>54.62</u>	28.46	<u>66.15</u>	24.62	53.33	<u>38.33</u>	67.69	<u>62.31</u>	59.23	<u>42.31</u>	<u>67.40</u>	<u>40.97</u>
Mixtral 8x22B	<u>80.77</u>	37.69	<u>72.31</u>	54.62	50.00	<u>42.62</u>	<u>37.69</u>	16.92	58.50	27.83	43.08	<u>44.62</u>	44.77	45.23	<u>55.63</u>	39.25
Deepseek I. 33B	70.77	37.23	68.46	46.31	53.08	<u>31.69</u>	51.54	37.69	50.00	27.50	46.92	26.15	<u>62.31</u>	39.23	57.58	35.11
LLAMA C. 34B	80.77	42.31	73.08	43.85	53.08	25.69	50.15	28.46	53.17	33.33	55.38	35.38	<u>46.92</u>	40.62	59.02	35.71
LLAMA-2 70B	78.46	24.62	71.54	36.77	53.08	26.92	51.38	16.15	60.83	22.00	49.23	55.54	24.46	26.00	55.72	29.71
LLAMA C. 70B	74.77	36.15	54.77	52.92	48.62	23.69	40.0	17.69	49.67	28.83	46.92	31.54	37.08	42.31	50.90	32.87
LLAMA-3 70B	90.77	82.31	93.08	86.15	87.69	82.31	78.62	<u>56.62</u>	60.50	<u>63.00</u>	62.31	<u>85.38</u>	78.15	64.77	78.71	74.30
LLAMA-3.1 70B	93.08	84.31	89.85	86.77	61.38	54.92	66.15	46.62	63.00	58.00	56.92	68.46	34.62	<u>69.23</u>	66.67	66.94
LLAMA-3.1 405B	<u>95.38</u>	86.92	<u>93.08</u>	93.85	59.23	<u>80.77</u>	<u>77.23</u>	62.92	65.00	<u>65.00</u>	90.00	86.92	<u>83.08</u>	65.38	<u>80.49</u>	77.42
GPT-4o Mini	90.77	73.85	95.38	79.23	80.77	39.23	67.69	46.15	54.17	21.67	77.69	70.00	76.92	67.69	77.74	56.50
GPT-4o	96.92	89.23	<u>94.62</u>	<u>90.00</u>	<u>79.23</u>	76.92	61.54	53.85	57.50	52.50	<u>88.46</u>	80.77	95.38	79.23	<u>81.84</u>	74.97

Table 2: Accuracy of 21 LLMs, (I)nstruct and (C)ode models, on 7 ACPBench tasks (boolean and multi-choice). The best results are **boldfaced**, second best are underlined, and the best among the small, open-sourced models are double underlined. All models were evaluated with two in-context examples and COT prompt. The right-most column is mean across tasks.



GPT-OSS 120B

Planning Benchmark Desiderata

- It should have a precise yet concise natural language description, including initial state, goal, and task dynamics.
- The problem should be sequential in nature, the order in which the actions need to be performed should matter.
- It should have a well defined action and state space.
- The problem should be of a non-trivial complexity.
- Must have sound validators for candidate solutions.
- It should have a large instance space and a dynamic generation procedure, thus allowing for the avoidance of memorization concerns.

Game of 24

Input: 4 numbers, Target: 24, Operations : [+, - , /, *]

Count Down

Input: set of x numbers, Target: y, Operations : [+, - , /, *]

Katz, Kokel and Sreedharan 2025

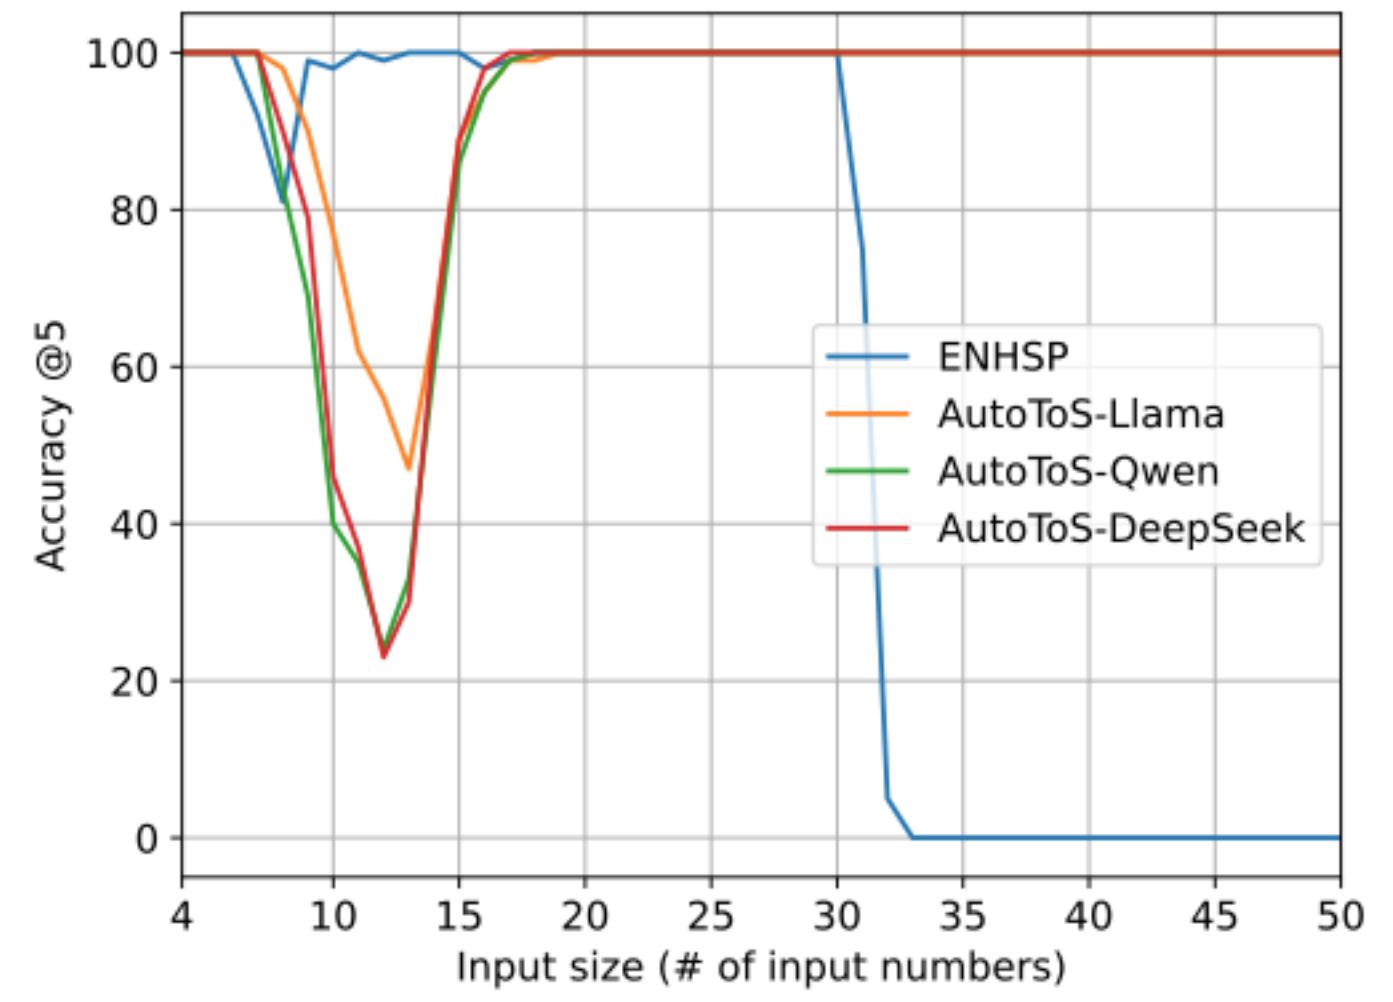


Figure 3: The accuracy of ENHSP and accuracy@5 of AutoToS with different language models for the Countdown problem.

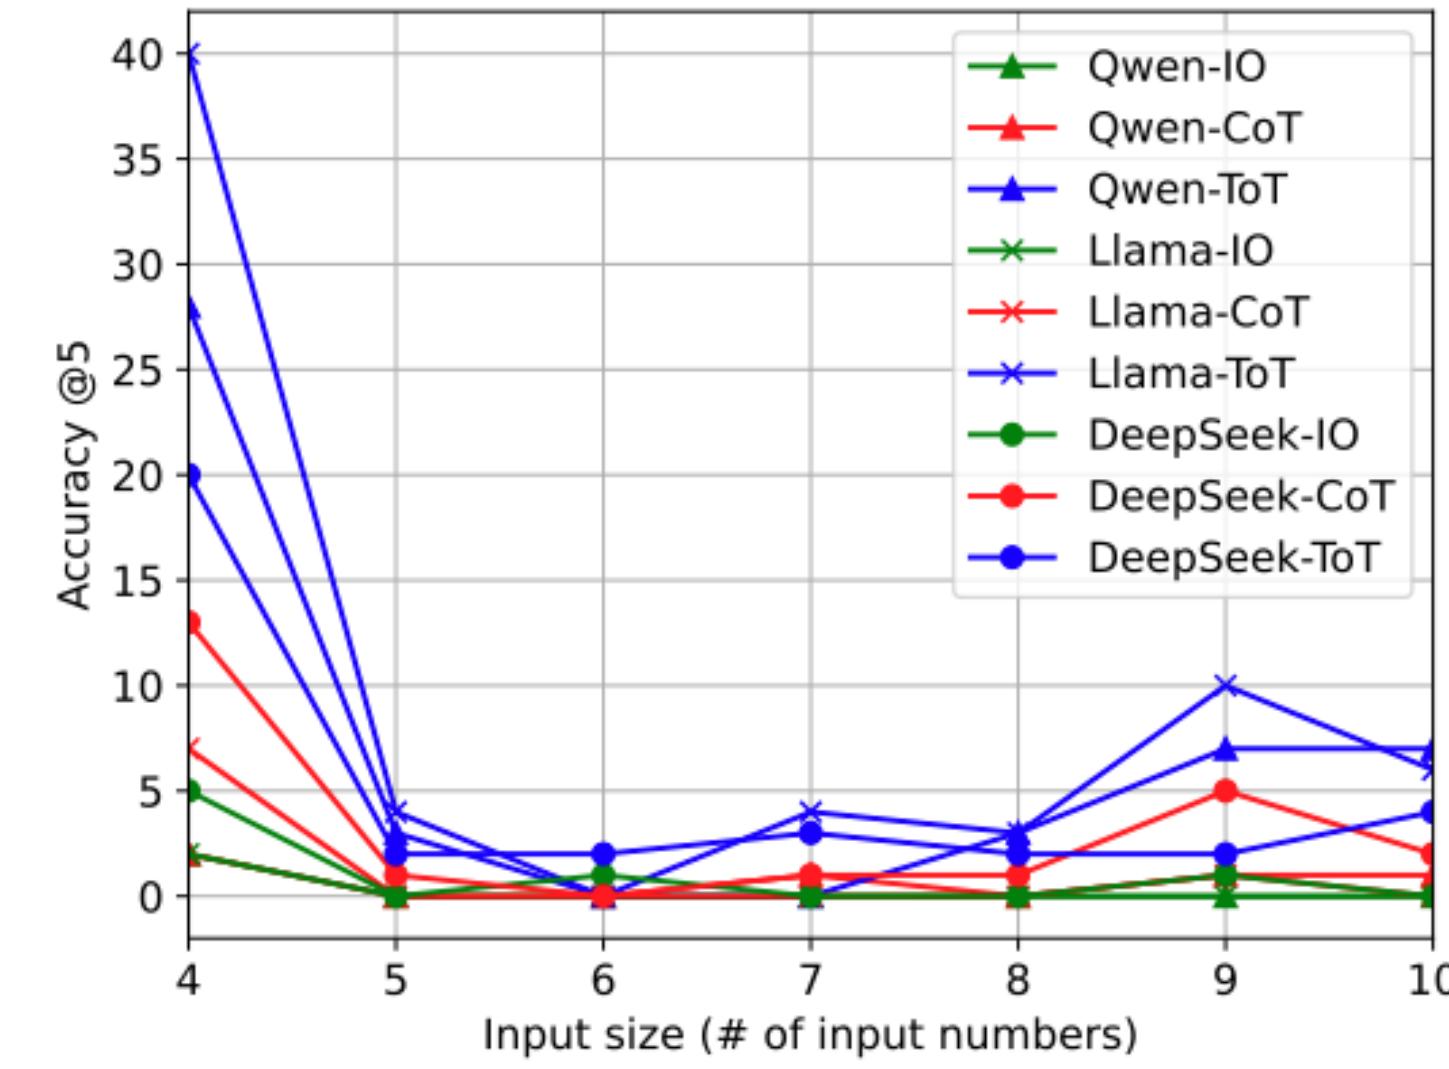


Figure 5: Accuracy @5 of LLM planning methods on CD.

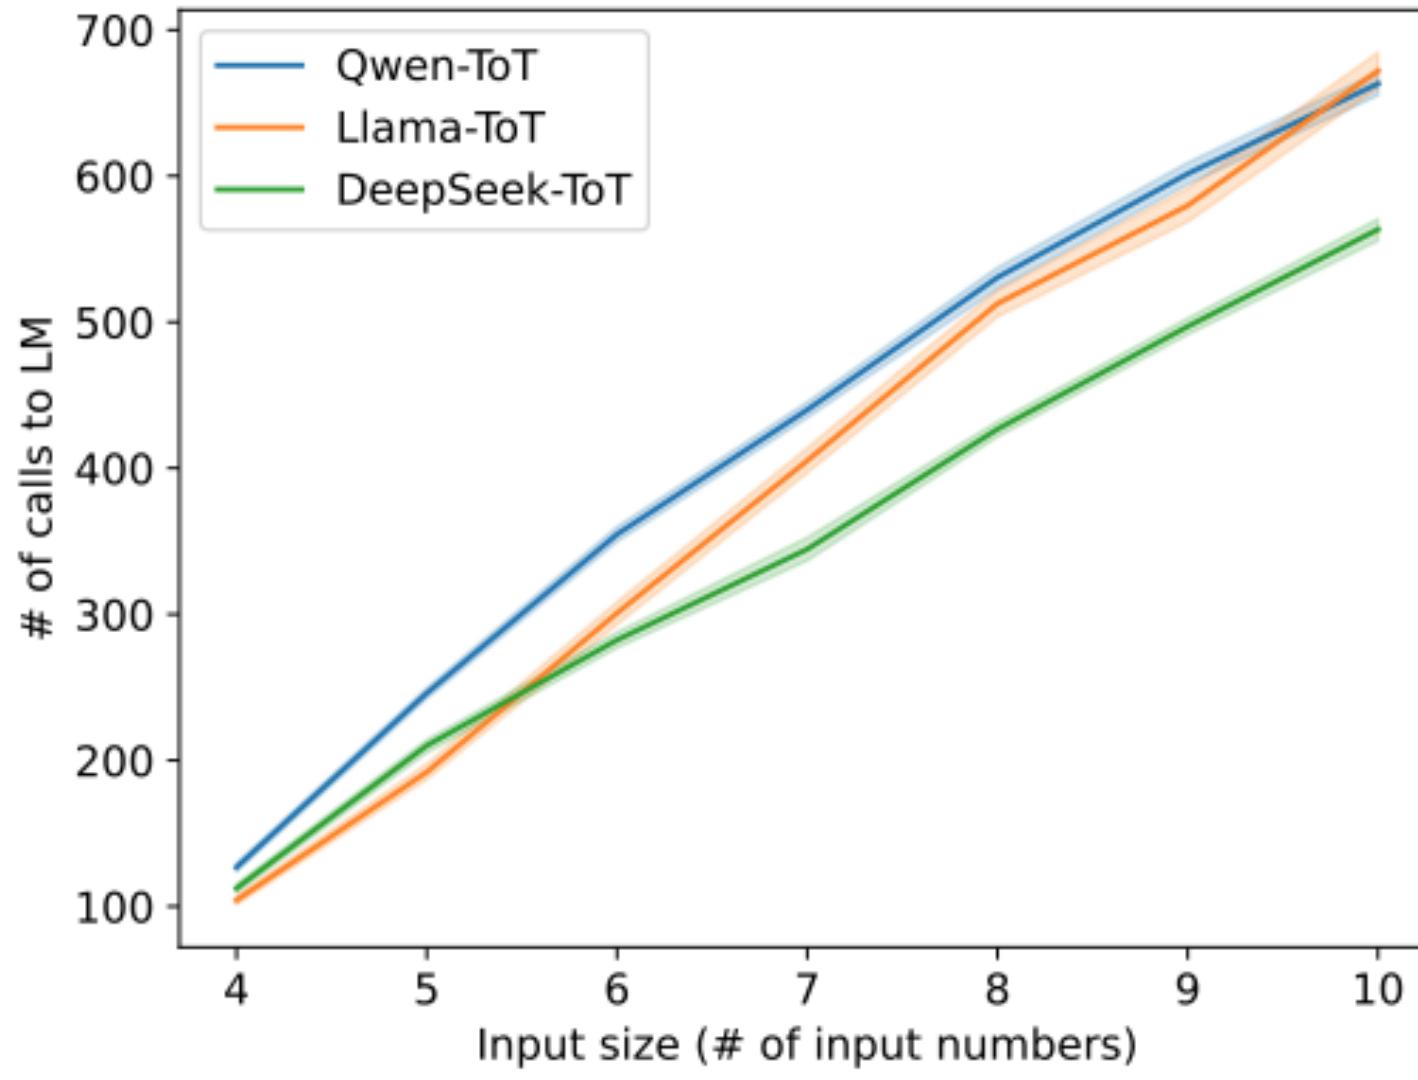


Figure 6: The average number of calls made to language models by the ToT approach with various language models.

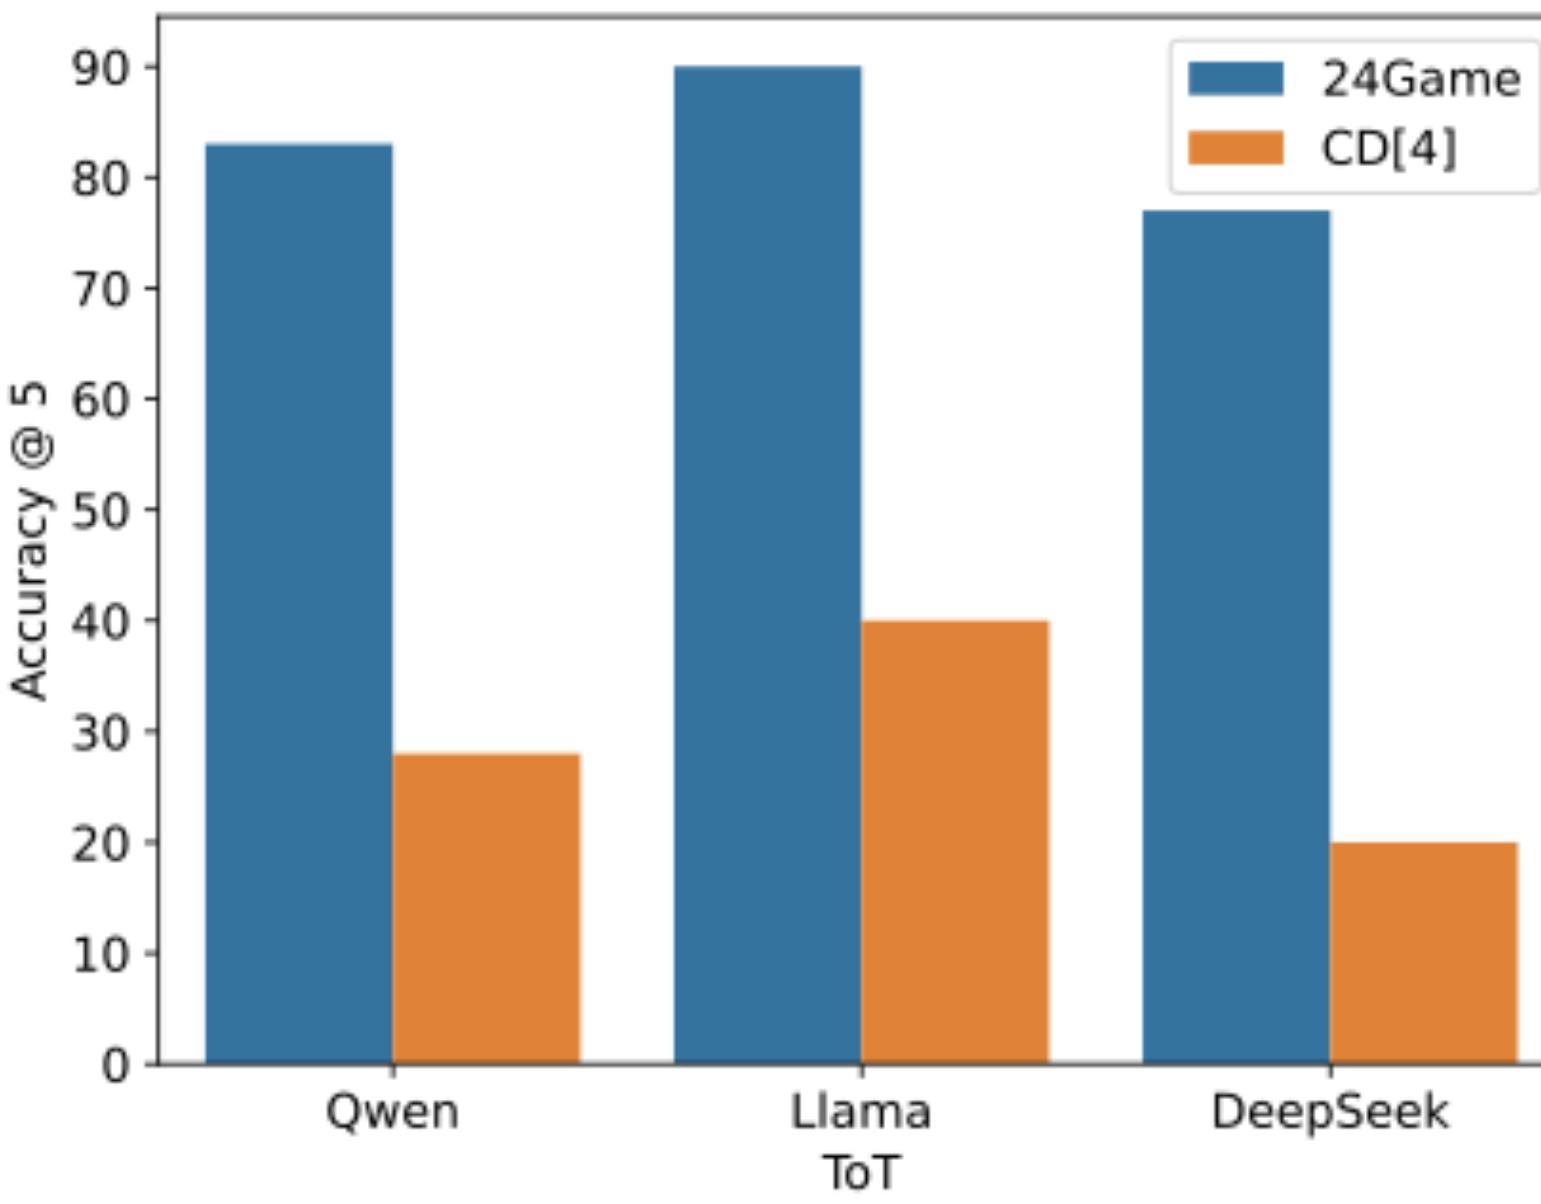
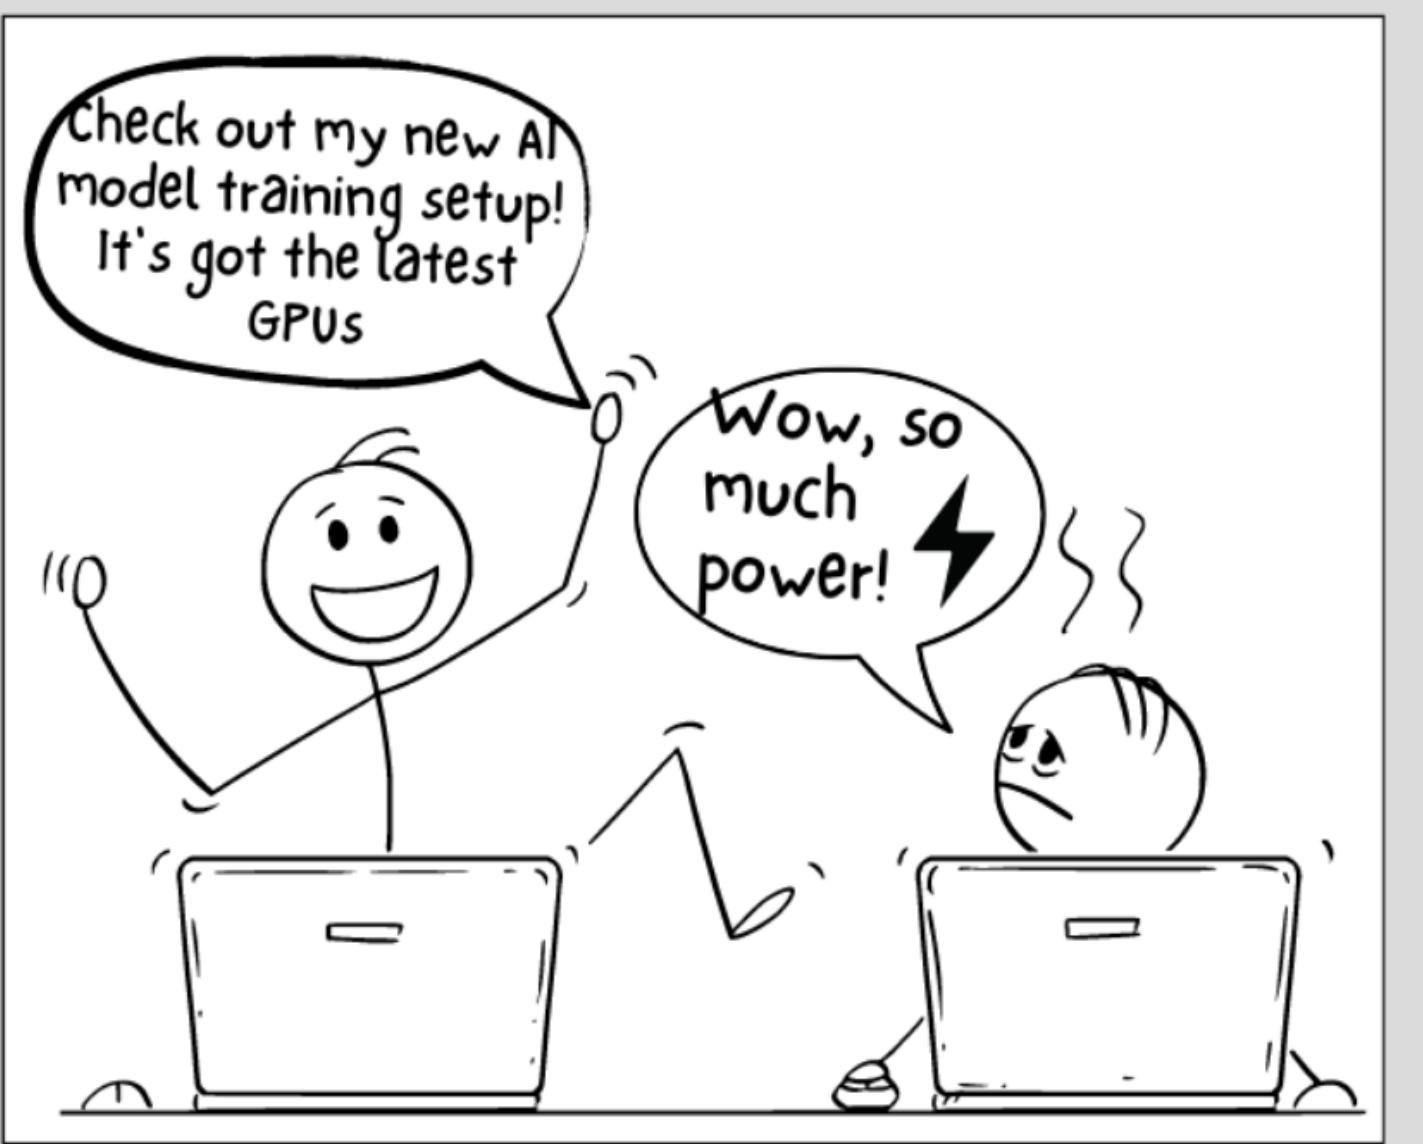


Figure 7: Accuracy @5 of various language models using the Tree of Thought (ToT) approach, comparing the 24Game dataset to instances of the same size (4) from our dataset.

Questions



The Carbon Impact of Large Language Models: AI's Growing Environmental Cost





IBM

≡

< Back to search results

RS: Toward Autonomous Data Management with FMs and AI Agents Intern 2026

Yorktown Heights, San Jose, Cambridge, Albany, New York, Massachusetts, California, United States IBM Research AI Internship

Email X LinkedIn

Apply now

Introduction

IBM Research takes responsibility for technology and its role in society. Working in IBM Research means you'll join a team who invent what's next in computing, always choosing the big, urgent and mind-bending work that endures and shapes generations. Our passion for discovery, and excitement for defining the future of tech, is what builds our strong culture around solving problems for clients and seeing the real world impact that you can make.

IBM's product and technology landscape includes Research, Software, and Infrastructure. Entering this domain positions you at the heart of IBM, where growth and innovation thrive.

Your role and responsibilities

We are looking for a talented and highly motivated intern to help advance our efforts in building autonomous data management systems. The work will include using foundation models (FMs) and AI agents for tasks such as data discovery, knowledge representation, data access and retrieval with querying, and automated data-driven insights. This is the core of the role: turning FMs and AI agents into dependable partners for real data work across modern data stacks.

Our focus is on research and development for data workflow orchestration — taking natural language all the way to trusted insights across multiple tools and functions. This includes step-by-step planning and reasoning for complex data workflows, developing low-computational-cost inference techniques so FMs and AI agents can efficiently automate or assist users on data tasks, and advancing trustworthy data agents where gauging factuality, faithfulness, and transparency of outputs over structured and unstructured data is critical. In this context, uncertainty quantification and mitigation, along with improved planning and reasoning, play a central role, providing useful tools to strengthen user confidence, manage model error propagation, and enable uncertainty-aware post-training. Together, these efforts make agentic solutions more efficient, accurate, and trustworthy.



LinkedIn: Harsha Kokel

References

- Katz, Michael, et al. "Thought of search: Planning with language models through the lens of efficiency." NeurIPS 2024
- Cao, Daniel, et al. "Automating Thought of Search: A Journey Towards Soundness and Completeness. (Student Abstract)" AAAI 2025
- Oswald, James, et al. "Large language models as planning domain generators." ICAPS 2024.
- Kokel, Harsha, et al. "ACPBench: Reasoning about Action, Change, and Planning." AAAI 2025.
- Kokel, Harsha, et al. "ACPBench Hard: Unrestrained Reasoning about Action, Change, and Planning." AAAI 2025 Workshop LM4Plan 2025.
- Wei, Jason, et al. "Chain-of-thought prompting elicits reasoning in large language models." NeurIPS 2022.
- Yao, Shunyu, et al. "Tree of thoughts: Deliberate problem solving with large language models." NeurIPS 2023.
- Yao, Shunyu, et al. "React: Synergizing reasoning and acting in language models." ICLR 2023.
- Xu, Binfeng, et al. "Rewoo: Decoupling reasoning from observations for efficient augmented language models." arXiv:2305.18323 2023.

References

- Kambhampati, Subbarao, et al. "Position: LLMs can't plan, but can help planning in LLM-modulo frameworks." ICML 2024.
- Huang, Wenlong, et al. "Language models as zero-shot planners: Extracting actionable knowledge for embodied agents." ICML 2022.
- Team, Gemini, et al. "Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context." arXiv:2403.05530 2024.
- Valmeekam, Karthik, et al. "LLMs Still Can't Plan; Can LRM? A Preliminary Evaluation of OpenAI's o1 on PlanBench." NeurIPS 2024 Workshop on Open-World Agents.
- Hao, Shibo, et al. "Reasoning with Language Model is Planning with World Model." EMNLP 2023.
- Besta, Maciej, et al. "Graph of thoughts: Solving elaborate problems with large language models." AAAI 2024.
- Zhou, Andy, et al. "Language Agent Tree Search Unifies Reasoning, Acting, and Planning in Language Models." ICML 2024.
- Liu, Bo, et al. "LLM+P: Empowering large language models with optimal planning proficiency." arXiv preprint arXiv:2304.11477 2023.

References

- He, Weinan, et al. "Exploring the capacity of pretrained language models for reasoning about actions and change." ACL 2023.
- Handa, Divij, et al. "ActionReasoningBench: Reasoning about Actions with and without Ramification Constraints." ICLR 2025.
- Guan, Lin, et al. "Leveraging pre-trained large language models to construct and utilize world models for model-based task planning." NeurIPS 2023